


# 2: Storage Latency


CPU

**C = \***



# Latencies


1. Read A
2. Read B
3.  $C = A * B$
4. Write C



With big data, most of the latency is memory latency (1,2,4), not computation (3)

# Storage Types

- Main Memory (RAM)



- Spinning disk



- Remote computer



# Summary for part 2

- The major source of latency in data analysis is reading and writing to storage
- Different types of storage offer different latency, capacity and price.
- Big data analytics revolves around methods for organizing storage and computation in ways that maximize speed while minimizing cost.
- Next, storage locality.

# 3: Caching


# Latency, size and price of computer memory

Given a budget, we need to trade off


\$10: Fast & Small

\$10: Slow & Large


# Cache: The basic idea




# Cache Hit




# Cache Miss




# Cache Miss Service: 1) Choose byte to drop



# Cache Miss Service: 2) write back



# Cache Miss Service: 3) Read In



# Summary of part 3

- Cache is much faster than main memory
- Cache hit = the needed value is already in the cache
- Cache miss = the needed value is not in the cache – needs to be brought in from memory
- If there is no space in cache:
  - need to make space
  - If dirty, need to write value back to memory first.
- Cache miss – latency is much bigger than cache miss.

# 4: Locality of storage access

# Access Locality

- The cache is effective If most accesses are hits.
  - Cache Hit Rate is high.
- **Cache effectiveness** depends on patterns (statistics) of memory access.
- **Temporal Locality**: Multiple accesses to **same** address within a short time period
- **Spatial Locality**: Multiple accesses to **near-by** addresses within a short time period

# Temporal Locality


- **Task:** compute the function  $f_{\theta}(x)$  on a long sequence  $x_1, x_2, \dots, x_n$
- $\theta$  is a parameter vector – example: the weights in a neural network.
- The parameters  $\theta$  are needed for each computation.
- If  $\theta$  fits in the cache – access is fast
- If  $\theta$  does **not** fit in the cache – each  $x_i$  causes at least one cache miss  
– program will be much slower.
- **Temporal Locality:** repeated access  
to the same memory location

# Spatial locality

- 
- **Task:** compute the function  $\sum_{i=1}^{n-1} (x_i - x_{i+1})^2$  on  $x_1, x_2, \dots, x_n$
- Contrast two ways to store  $x_1, x_2, \dots, x_n$  :
- Linked list (poor locality)
- Indexed array (good locality)

# Linked List

Let  $x_1, x_2, \dots, x_n$  be 1,2,3,4,5,6



Traversal of 6 elements touches 4 pages

# array

Let  $x_1, x_2, \dots, x_n$  be 1,2,3,4,5,6

Page 1	Page 2			Page3			Page4			
	1	2	3	4	5	6				

Traversal of 6 elements touches 2 pages

# Summary of Part 4

- Caching Is effective when memory access is local
- Temporal locality: accessing the same location many times in a short period of time.
- Spatial locality: accessing close-by locations many times in a short period of time.
- Hardware and compilers have a symbiotic relationship: success if compiler generates machine code that has good locality.

# Word Count

- Task: given a (large) text
- Count the number of times each word
- Output (word,count) sorted in decreasing order by Count.

# Unsorted word count / poor locality

```
==== unsorted list:  
the, vernacular, but, as, for, you, ye, carrion, rogues, turning, to,
```

Dict={}

For word in list:

    if word in Dict:

        Dict[word]+=1

    else:

        Dict[word]=1

Suppose

    len(list)=1,000,000

    len(Dict) = 100,000

Access to list: spatially local

Access to Dict: **random**

# sorted word count / good locality

==== sorted list:

lines, lingered, lingered, lingered, lingered, lingered, lingered, lingerin  
g, lingering, lingering, lingering, lingering, lingering, lingeri  
ng, lingering, lingers, lingo, lingo, lining, link, link, linked, li  
nk, linked, linked, links, links

Dict={}

Sort(list)

For word in list:

    if word in Dict:

        Dict[word]+=1

    else:

        Dict[word]=1

Suppose

len(list)=1,000,000

len(Dict) = 100,000

Access to list: spatially local

Access to Dict: **Spatially local**

**But what about the sort step?**

Sorting can be done in time **O(n)**

Efficient in distributed setup

# Summary

- Improved memory locality reduces run-time
- Why?
  - Because computer memory is organized in pages.
  - And caching retrieves a page at a time.