

DSC 232R: Big Data Analytics Using Spark

Winter 2026

Week 2

January 18, 2026

1 Topic: HDFS and The MapReduce Paradigm

1.1 The Distributed File System (HDFS)

1.1.1 Lecture Content

- **The Core Problem:** Standard file systems cannot handle petabytes of data because they rely on a single disk/computer.
- **Google's Solution (GFS/HDFS):**
 1. **Commodity Hardware:** Use thousands of cheap, unreliable computers instead of one expensive supercomputer.
 2. **Chunking:** Files are broken into fixed-size blocks (default 128MB) called "Chunks" or "Splits".
 3. **Redundancy (Replication):** Every chunk is copied to multiple machines (Default Replication Factor = 3).
 - *Why?* If a machine crashes (which happens constantly), the data exists elsewhere.
 4. **Architecture:**
 - **NameNode (Master):** Stores metadata (filename → list of chunk IDs).
 - **DataNode (Worker):** Stores the actual raw bytes.
- **Locality Principle:** "Move computation to the data, not data to the computation."
 - Network bandwidth is the bottleneck. It is faster to send a 5KB program to the data than to pull 1TB of data to the program.

2 Topic: Spark Architecture

2.1 RDDs and Context

2.1.1 Lecture Content

- **SparkContext (The Driver):**

- The "Brain" of the operation. Resides on the Master node.
- Does NOT store the data. It stores the *Execution Plan* (Lineage).
- Coordinates workers (Executors).

- **RDD (Resilient Distributed Dataset):**

- **Resilient:** Reconstructs lost data automatically using lineage (re-computing steps) rather than replication.
- **Distributed:** Partitioned across the cluster.
- **Immutable:** Once created, it cannot be changed. You can only create a *new* RDD from an old one.

- **Spark vs. Hadoop:**

- **Hadoop MapReduce:** writes to Disk after every step (High Latency).
- **Spark:** keeps data in Memory (RAM) between steps (Low Latency).

2.2 Lazy Evaluation and Pipelines

2.2.1 Lecture Content

- **Two Types of Operations:**

1. **Transformations (Lazy):** Definition of a plan. Returns a new RDD.
 - *Examples:* `map`, `filter`, `flatMap`, `sample`.
 - *Behavior:* Spark does NOTHING when these lines are run. It builds a DAG (Directed Acyclic Graph).
2. **Actions (Eager):** Triggers execution. Returns a result to the Driver.
 - *Examples:* `collect`, `count`, `reduce`, `take`, `saveAsTextFile`.
 - *Behavior:* Spark optimizes the DAG and launches tasks on the cluster.

- **Pipelining (Fusion):**

- Because of lazy evaluation, Spark can combine multiple steps into one pass.

- *Example:* `map` + `filter` happens in a single loop over the data, avoiding intermediate memory storage.

- **Caching:**

- Because RDDs are re-computed by default, branching (using an RDD twice) causes double computation.
- `rdd.cache()` tells Spark to save the result in RAM after the first computation.

3 Topic: Key-Value Operations (The Danger Zone)

3.1 Pair RDDs

3.1.1 Jupyter Notebook Content

- **Data Structure:** RDDs where elements are Python tuples: `(Key, Value)`.
- **reduceByKey(func):**
 - Groups values by key and applies `func` (e.g., sum).
 - **Architecture Win:** Performs *Map-Side Combination*. It reduces data LOCALLY before sending it over the network.
 - *Result:* Low network traffic, high performance.
- **groupByKey():**
 - Groups values by key but performs NO reduction. Returns `(Key, Iterable)`.
 - **Architecture Fail:** Sends ALL raw data over the network to the reducer.
 - *Result:* High network traffic, likely to cause **Out Of Memory (OOM)** errors on the reducer.
- **The Shuffle:**
 - Operations that require moving data between partitions (by key) cause a "Shuffle".
 - This breaks the pipeline and defines a "Stage Boundary".
 - *Shuffle Ops:* `reduceByKey`, `groupByKey`, `join`, `repartition`.

4 Exam Traps: Danger Zone

- **The "Average" Trap (Algebraic Properties):**
 - You cannot calculate an average using `reduce(lambda x,y: (x+y)/2)`.
 - **Why?** Reduce operations must be **Associative** and **Commutative**.
 - $(a + b)/2 + c$ is not the same as $(a + b + c)/3$.
 - **Solution:** Map to `(sum, count)` tuples first, reduce by adding components element-wise, then divide at the very end.
- **The groupByKey Trap:**
 - Prof will ask: "Why did my cluster crash when I switched from `reduceByKey` to `groupByKey`?"

- **Answer:** `groupByKey` forces a massive shuffle of all data. `reduceByKey` pre-aggregates on the map side (like a combiner), sending minimal data.

- **The Collect Trap:**

- `rdd.collect()` brings ALL data to the Driver node.
- If the RDD is 1TB and the Driver has 16GB RAM, the driver crashes.
- **Exam Fix:** Use `take(n)` or `saveAsTextFile()` instead.

- **Terminology Shift:**

- Lecture: "Lazy Evaluation". Exam: "DAG Construction vs. Materialization".
- Lecture: "Shuffle". Exam: "Wide Dependency".
- Lecture: "Chunk". Exam: "Partition" or "Split".

5 Practice Quiz

Question 1

The Philosophy of MapReduce

Which statement best describes the "Data Locality" principle in HDFS/Spark?

- a All data is gathered to the Master node for processing.
- b Data is randomly shuffled across the network to ensure fairness.
- c Code (computation) is sent to the node where the data resides.
- d Data is always compressed before processing.

Answer: C

Brief Explanation

- Network bandwidth is the scarcest resource. Moving code (kilobytes) to the data (terabytes) is orders of magnitude faster than moving data to the code.

Question 2

The Math of Reduce

Why does the operation `rdd.reduce(lambda a,b: a - b)` yield unpredictable results?

- a Subtraction is too computationally expensive.
- b Subtraction is not Associative.
- c Spark does not support negative numbers.
- d The lambda syntax is incorrect.

Answer: B

Brief Explanation

- Reduce operations execute in parallel in undefined orders. $(a - b) - c$ is not the same as $a - (b - c)$. Operations must be Associative and Commutative to guarantee deterministic results in a distributed system.

Question 3

Lazy Evaluation

You run the following code on a 10TB dataset:

```
data = sc.textFile("huge_file.txt")
mapped = data.map(lambda x: x.split())
filtered = mapped.filter(lambda x: len(x) > 5)
```

How much time does this take to execute?

- a Approx 1 hour (reading 10TB).
- b Approx 10 minutes (filtering reduces size).
- c Microseconds (almost instant).
- d It depends on the number of executors.

Answer: C

Brief Explanation

- `textFile`, `map`, and `filter` are all **Transformations**. Spark is Lazy—it only builds the Execution Plan (DAG) in memory. No data is read until an Action (like `count`) is called.

Question 4

Performance Tuning

You need to count word frequencies. Which operation is more efficient and why?

- a `groupByKey().map(sum)`
- b `reduceByKey(sum)`

Answer: B

Brief Explanation

- `reduceByKey` performs **Map-Side Combination**. It sums the counts locally on each worker *before* sending data across the network. `groupByKey` shuffles every single word pair, causing massive network congestion.

Question 5

Stage Boundaries

What triggers the creation of a new "Stage" in a Spark Job?

- a Any Transformation.

- b Any Action.
- c A Shuffle (Wide Dependency).
- d A Cache command.

Answer: C

Brief Explanation

- Spark pipelines "Narrow Dependencies" (map, filter) into a single stage. A "Shuffle" (reduceByKey, join, repartition) requires data to move between partitions, forcing a hard barrier (Stage Boundary) where all previous tasks must finish before the next begins.

Question 6

HDFS Reliability

If a DataNode crashes, how does HDFS ensure data is not lost?

- a It uses RAID 5 on the Master Node.
- b It recovers the data from the NameNode's RAM.
- c It relies on the 3x Replication Factor (copies exist on other nodes).
- d The job simply fails.

Answer: C

Brief Explanation

- HDFS replicates every chunk (default 3 times) across different nodes. If one node dies, the Master simply redirects requests to one of the replicas.

Question 7

Driver vs. Executor

What is the primary risk of running the following command on a production cluster? `result = huge_rdd.collect()`

- a It will delete the original data on HDFS.
- b It forces the Executors to crash.
- c It causes an Out Of Memory (OOM) error on the Driver.
- d It runs too slowly because of Python serialization.

Answer: C

Brief Explanation

- `collect()` fetches all data from all distributed partitions and tries to fit it into the memory of the single Driver machine. If the data *>* Driver RAM, the driver crashes.

Question 8

DAG Visualization

What is the purpose of `rdd.toDebugString()`?

- a To print the first 10 rows of data.
- b To view the logical Execution Plan (Lineage).

- c To check for syntax errors in Python.
- d To check the size of the RDD in bytes.

Answer: B

Brief Explanation

- `toDebugString()` outputs the DAG (Directed Acyclic Graph) showing the lineage of transformations and where shuffles/caching will occur.

Question 9

Glom

What does the `glom()` transformation do?

- a Deletes empty partitions.
- b Coalesces all partitions into one.
- c Transforms each partition into a list (array) of elements.
- d Sorts the data globally.

Answer: C

Brief Explanation

- `glom()` maps each partition to a single array containing all elements of that partition. It is useful for efficient batch operations or checking partition skew (e.g., `glom().map(len).collect()`).

Question 10

Immutability

You want to add 5 to every element in an RDD ‘A’. You run ‘`A.map(lambda x: x+5)`’. What happens to ‘A’?

- a ‘A’ is modified in place.
- b ‘A’ is deleted and replaced by the new result.
- c ‘A’ remains unchanged; a new RDD is returned.
- d It depends on whether ‘A’ is cached.

Answer: C

Brief Explanation

- RDDs are **Immutable**. You cannot change an RDD. Transformations always return a *pointer* to a new RDD that depends on the old one.