

Dataframe operations

Spark DataFrames allow operations similar to Pandas dataframes. We demonstrate some of those.

For more see the official guide and this article.

But first, some important details

Spark and datahub can run in two modes:

- **Local mode** which means that it is run on the same computer as the head node. This convenient for small jobs and for debugging. Can also be done on your laptop.
- **Remote mode** in which the head node and the worker nodes are separate. This mode requires that the spark cluster is up and running. In this case the full resources of the clusters are available. This mode is useful when processing hard jobs.

In [1]:

```
import os
import sys

import pyspark
from pyspark import SparkContext
from lib import sparkConfig
```

172.17.0.2

In [2]:

```
%%time
#sc.stop() # uncomment if sparkContext already exists
sc = SparkContext() #conf=sparkConfig.conf)
sc
```

```
CPU times: user 15.7 ms, sys: 5.82 ms, total: 21.6 ms
Wall time: 1.58 s
```

Out[2]:

SparkContext

[Spark UI](#)

Version

v3.2.1

Master

local[*]

AppName

pyspark-shell

In [3]:

```
%%time
from pyspark import SparkContext
from pyspark.sql import SQLContext
from pyspark.sql.types import Row, StructField, StructType, StringType,
                               BooleanType, IntegerType, DoubleType, DateType
sqlContext = SQLContext(sc)
sqlContext
```

```
CPU times: user 2.24 ms, sys: 1.84 ms, total: 4.08 ms
Wall time: 69.7 ms
```

```
/usr/local/spark/python/pyspark/sql/context.py:77: FutureWarning
  Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate()
  instead.
  warnings.warn(
```

Out[3]: <pyspark.sql.context.SQLContext at 0xfffffacbed190>

In [4]:

```
%%time  
parquet_path='/datasets/weather/datasets/weather-parquet/'  
df=sqlContext.read.parquet(parquet_path)
```

```
CPU times: user 2.93 ms, sys: 1.11 ms, total: 4.04 ms  
Wall time: 1.76 s
```

In [4]:

```
%time  
parquet_path='/datasets/weather/datasets/weather-parquet/'  
df=sqlContext.read.parquet(parquet_path)
```

```
CPU times: user 2.93 ms, sys: 1.11 ms, total: 4.04 ms  
Wall time: 1.76 s
```

In [5]:

```
df.printSchema()
```

```
root  
|-- Station: string (nullable = true)  
|-- Measurement: string (nullable = true)  
|-- Year: integer (nullable = true)  
|-- Values: binary (nullable = true)
```

In [6]:

```
%%time  
print(df.count())  
df.show(1)
```

12720632

Station	Measurement	Year	Values
AG000060390	TAVG	2022	[79 00 6C 00 66 0...]

only showing top 1 row

CPU times: user 0 ns, sys: 2.91 ms, total: 2.91 ms

Wall time: 1.58 s

`.(describe)`

The method `df.describe()` computes five statistics for each column of the dataframe `df`.

The statistics are **count,mean,std,min,max**

```
In [7]: df.describe().select('Station','Measurement','Year').show()
```

Station	Measurement	Year
12720632	12720632	12720632
null	null	1977.4126412901496
null	null	31.931757686816614
ACW00011604	ACMC	1764
ZI000067991	WV20	2022

groupby and agg

The method `groupby(col)` groups rows according the value of the column `(col)`.

The method `.agg(spec)` computes a summary for each group specified in `spec`

```
In [8]: df.groupby('Measurement').agg({'Year': 'min', 'Station': 'count'}).shc
```

Measurement	count(Station)	min(Year)
WESD	123619	1952
PGTM	31995	1948
AWDR	730	1996
SX13	60	1982
WT07	19790	1891
WT13	7659	1996
SX33	524	1982
EVAP	27035	1893
SN53	201	1982
WT10	2487	1886
SN35	59	2005
TMIN	1388531	1764
WT21	2434	1996
DATX	7766	1863
WT14	38509	1851
WT19	4928	1992
SX02	3373	1982
MDPR	515120	1832
WT09	30516	1852
SX51	178	1982

only showing top 20 rows

Using SQL Queries on DataFrames

There are two main ways to manipulate DataFrames:

Imperative Manipulation

Using python methods such as `.select` and `.groupby`.

- Advantage: order of operations is specified
- Disadvantage: you need to describe both **what** is the result you want and **how** to get it.

Declarative Manipulation (SQL)

- Advantage: You need to describe only **what** is the result you want
- Disadvantage: SQL does not have primitives for common analysis operations such as **covariance**

Counting the number of occurrences of each measurement,
comparatively

In [10]:

```
%time  
L=df.groupBy('measurement').count().collect()  
#L is a list of Rows (collected DataFrame)
```

CPU times: user 1.98 ms, sys: 2.28 ms, total: 4.26 ms
Wall time: 1.02 s

```
In [10]:
```

```
%time  
L=df.groupBy('measurement').count().collect()  
#L is a list of Rows (collected DataFrame)
```

```
CPU times: user 1.98 ms, sys: 2.28 ms, total: 4.26 ms  
Wall time: 1.02 s
```

```
In [11]:
```

```
D=[(e['measurement'],e['count']) for e in L]  
print('The most common mesurements')  
sorted(D,key=lambda x:x[1], reverse=True)[:6]
```

```
The most common mesurements
```

```
Out[11]:
```

```
[('PRCP', 3256604),  
 ('TMIN', 1388531),  
 ('TMAX', 1384761),  
 ('SNOW', 1257297),  
 ('SNWD', 1208442),  
 ('TOBS', 536759)]
```

```
In [12]: print('The most rare mesurements')
sorted(D,key=lambda x:x[1], reverse=False)[:6]
```

The most rare mesurements

```
Out[12]: [('SN57', 1), ('SX15', 2), ('SX57', 2), ('SX17', 2), ('SN14',
2), ('SX14', 3)]
```

Counting the number of occurrences of each measurement,
declaratively

Registering a dataframe as a table in a database.

In order to apply SQL commands to a dataframe, it has to first be registered as a table in the database managed by sqlContext.

```
In [13]: sqlContext.registerDataFrameAsTable(df, 'weather') #using older sqlConte
```

In [14]:

```
%%time
query"""
SELECT measurement,COUNT(measurement) AS count,
       MIN(year) AS MinYear
FROM weather
GROUP BY measurement
ORDER BY count DESC
"""
print(query)
sqlContext.sql(query).show(5)
```

```
SELECT measurement,COUNT(measurement) AS count,
       MIN(year) AS MinYear
```

```
FROM weather
GROUP BY measurement
ORDER BY count DESC
```

measurement	count	MinYear
PRCP	3256604	1781
TMIN	1388531	1764
TMAX	1384761	1764
SNOW	1257297	1840
SNWD	1208442	1857

only showing top 5 rows

CPU times: user 2.53 ms, sys: 1.2 ms, total: 3.73 ms

Performing a map command

- Dataframes do not support `map` and `reduce` operations.
- In order to perform a `map` or `reduce` on a dataframe, you first need to transform it into an RDD

Performing a map command

- Dataframes do not support `map` and `reduce` operations.
- In order to perform a `map` or `reduce` on a dataframe, you first need to transform it into an RDD
- This is a quick-and-dirty solution.
- A better way is to use [built-in sparkSQL functions](#)
- Or if you can't find what you need, you can try and create a [User-Defined-Function*](#) (UDF)

```
In [15]: df.rdd.map(lambda row:(row.Station,row.Year)).take(5)
```

```
Out[15]: [('AG000060390', 2022),  
          ('AGE00147716', 2022),  
          ('AGM00060360', 2022),  
          ('AGM00060421', 2022),  
          ('AGM00060430', 2022)]
```

Aggregations

- **Aggregation** can be used, in combination with built-in spark SQL functions, to compute statistics of a dataframe.
- Computation will be fast thanks to combined optimizations with database operations.
- A partial list: `count()`, `approx_count_distinct()`, `avg()`,
`max()`, `min()`
- Of these, the interesting one is `approx_count_distinct()` which uses sampling to get an approximate count fast.

```
In [17]: df.agg({'station':'approx_count_distinct'}).show()
```

approx_count_distinct(station)
128546

Approximate Quantile

- Suppose we want to partition the years into 10 ranges
- Such that in each range we have approximately the same number of records.
- The method `approxQuantile` will use a sample to do this for us

Approximate Quantile

- Suppose we want to partition the years into 10 ranges
- Such that in each range we have approximately the same number of records.
- The method `approxQuantile` will use a sample to do this for us

```
In [27]: %time
print('with accuracy 0.1: ',df.approxQuantile('Year', [0.1*i for i in r
with accuracy 0.1:  [1764.0, 1948.0, 1961.0, 1973.0, 1980.0, 19
90.0, 1997.0, 2006.0, 2022.0]
CPU times: user 7.28 ms, sys: 2.84 ms, total: 10.1 ms
Wall time: 947 ms
```

Approximate Quantile

- Suppose we want to partition the years into 10 ranges
- Such that in each range we have approximately the same number of records.
- The method `approxQuantile` will use a sample to do this for us

In [27]:

```
%%time
print('with accuracy 0.1: ',df.approxQuantile('Year', [0.1*i for i in r
with accuracy 0.1:  [1764.0, 1948.0, 1961.0, 1973.0, 1980.0, 19
90.0, 1997.0, 2006.0, 2022.0]
CPU times: user 7.28 ms, sys: 2.84 ms, total: 10.1 ms
Wall time: 947 ms
```

In [28]:

```
%%time
print('with accuracy 0.001: ',df.approxQuantile('Year', [0.1*i for i in r
with accuracy 0.001:  [1931.0, 1951.0, 1962.0, 1972.0, 1982.0,
1991.0, 2000.0, 2009.0, 2015.0]
CPU times: user 11.5 ms, sys: 6.71 ms, total: 18.2 ms
Wall time: 5.19 s
```

Reading rows selectively from Parquet

Suppose we are only interested in snow measurements. We can apply an SQL query directly to the Parquet files. As the data is organized in columnar structures, we can do the selection efficiently without loading the whole file to memory.

Here the file is small, but in real applications it can consist of hundreds of millions of records. In such cases loading the data first to memory and then filtering it is very wasteful.

In [29]:

```
query="""SELECT station,measurement,year
FROM parquet.`ts`
WHERE measurement=\"SNOW\" """%parquet_path
print(query)
df2 = sqlContext.sql(query)
print(df2.count(),df2.columns)
df2.show(5)
```

```
SELECT station,measurement,year
FROM parquet.`/datasets/weather/datasets/weather-parquet/` 
WHERE measurement="SNOW"
1257297 ['station', 'measurement', 'year']
+-----+-----+----+
|    station|measurement|year|
+-----+-----+----+
|BF1SS000005|      SNOW|2022|
|CA001016335|      SNOW|2022|
|CA0010253G0|      SNOW|2022|
|CA00102BFHH|      SNOW|2022|
|CA001036570|      SNOW|2022|
+-----+-----+----+
only showing top 5 rows
```

Summary

- Dataframes can be manipulated declaratively, which allows for more optimization.
- Dataframes can be stored and retrieved from Parquet files.
- It is possible refer to directly to a parquet file in an SQL query.
- See you next time!