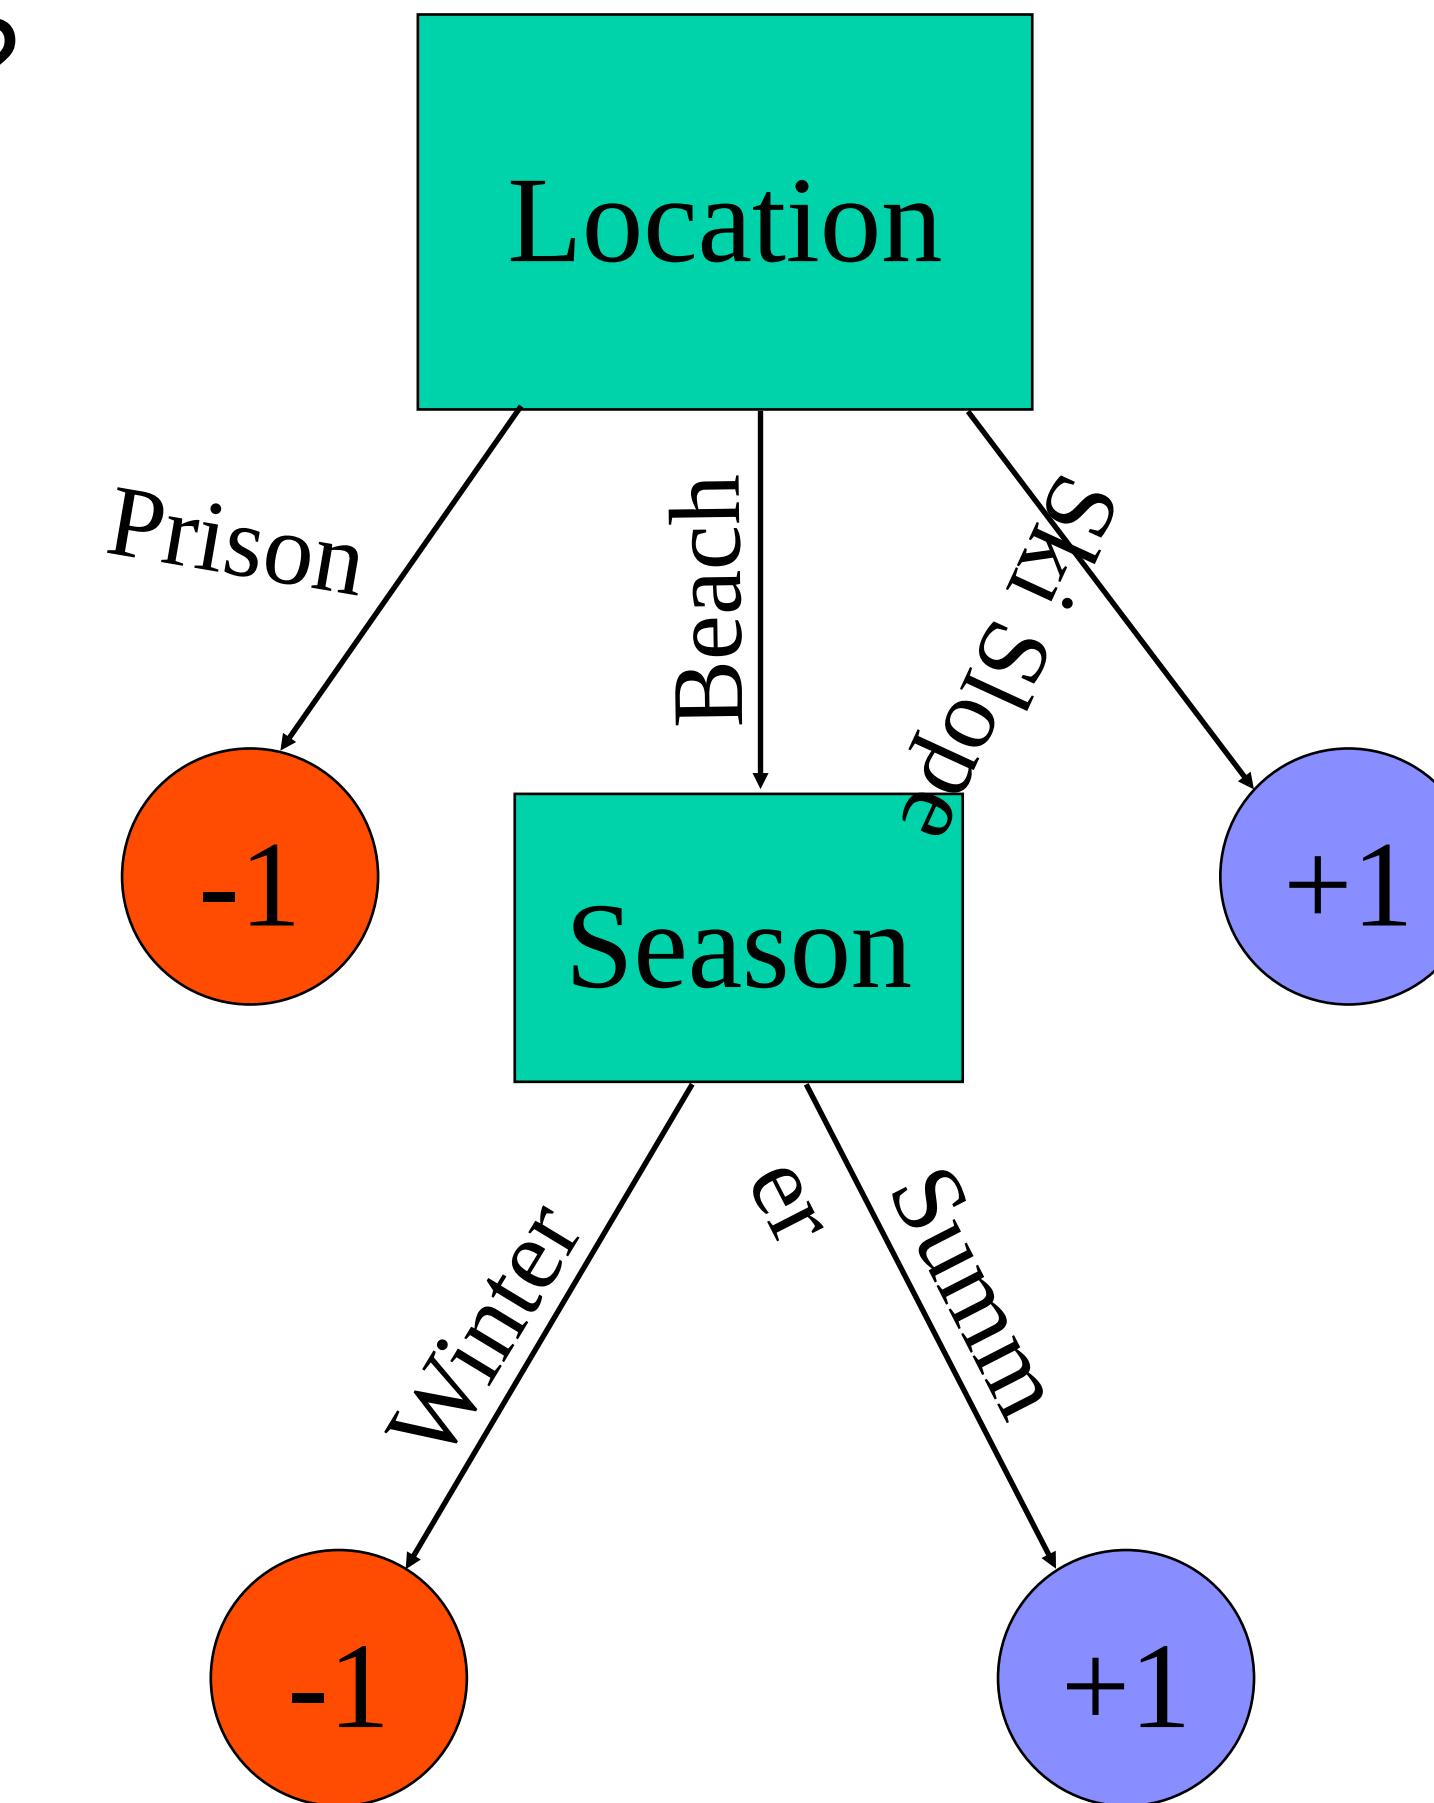


# Decision trees

# Decision Trees / Discrete Features

Where and when do you have fun?

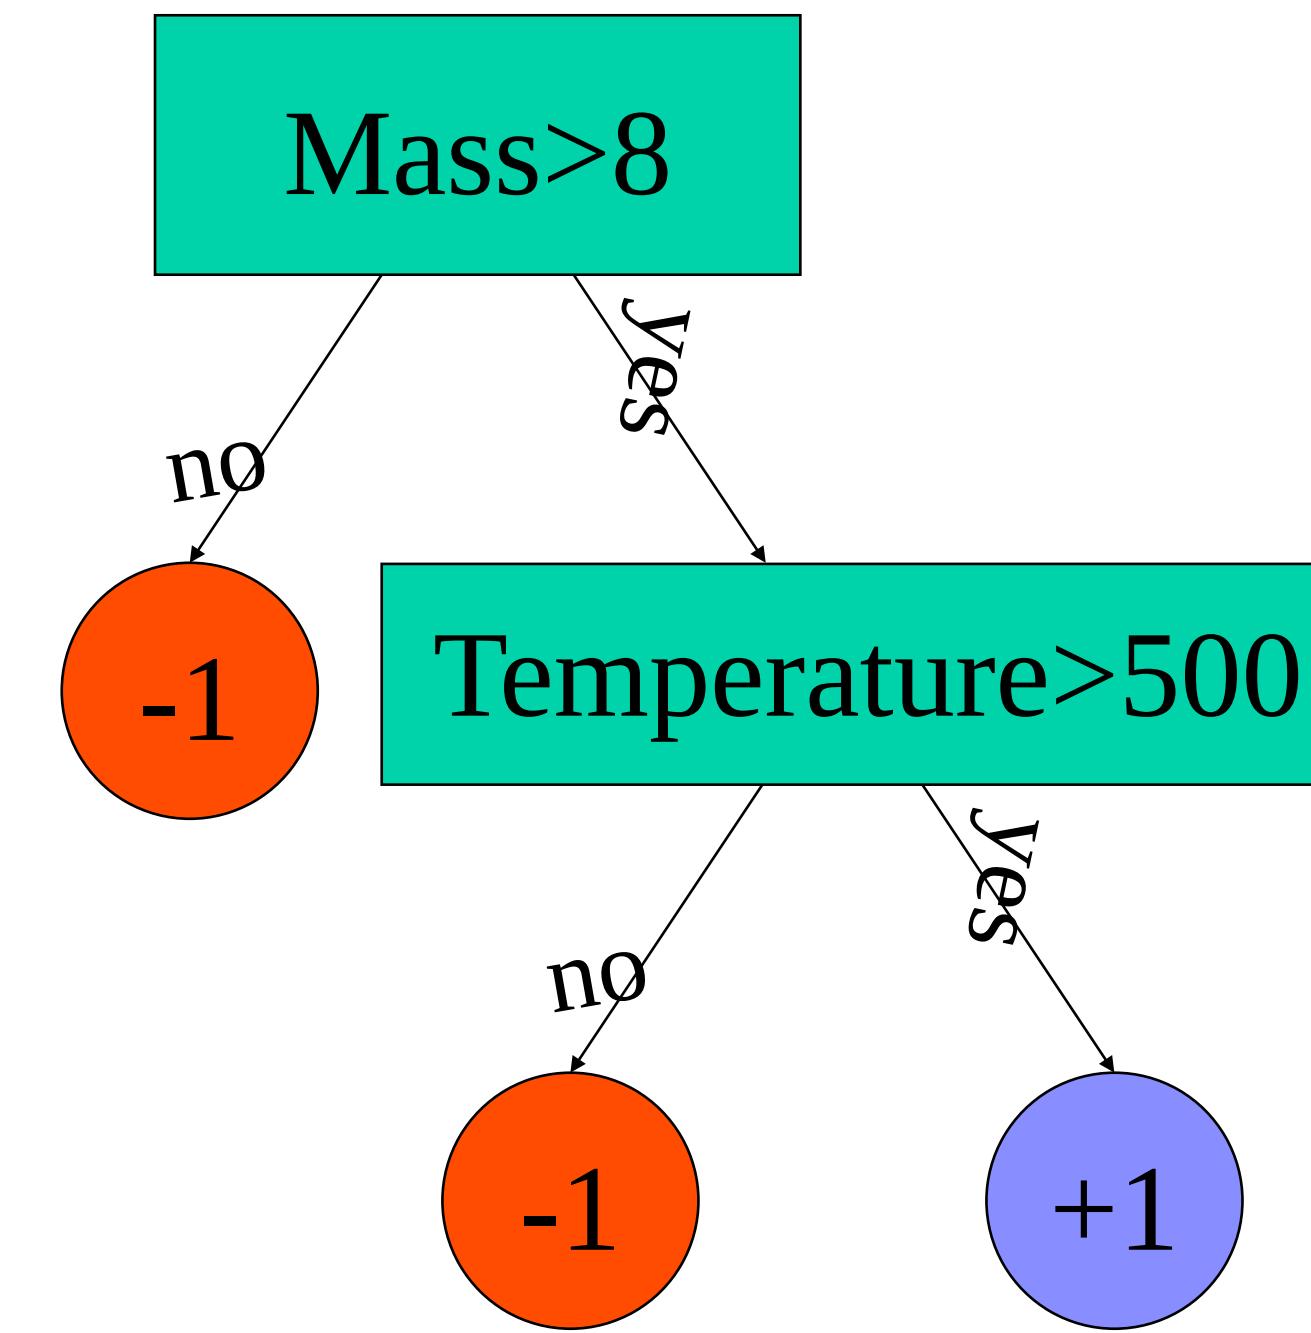
Season	Location	Fun?
summer	prison	-1
summer	beach	+1
Winter	ski-slope	+1
Winter	beach	-1



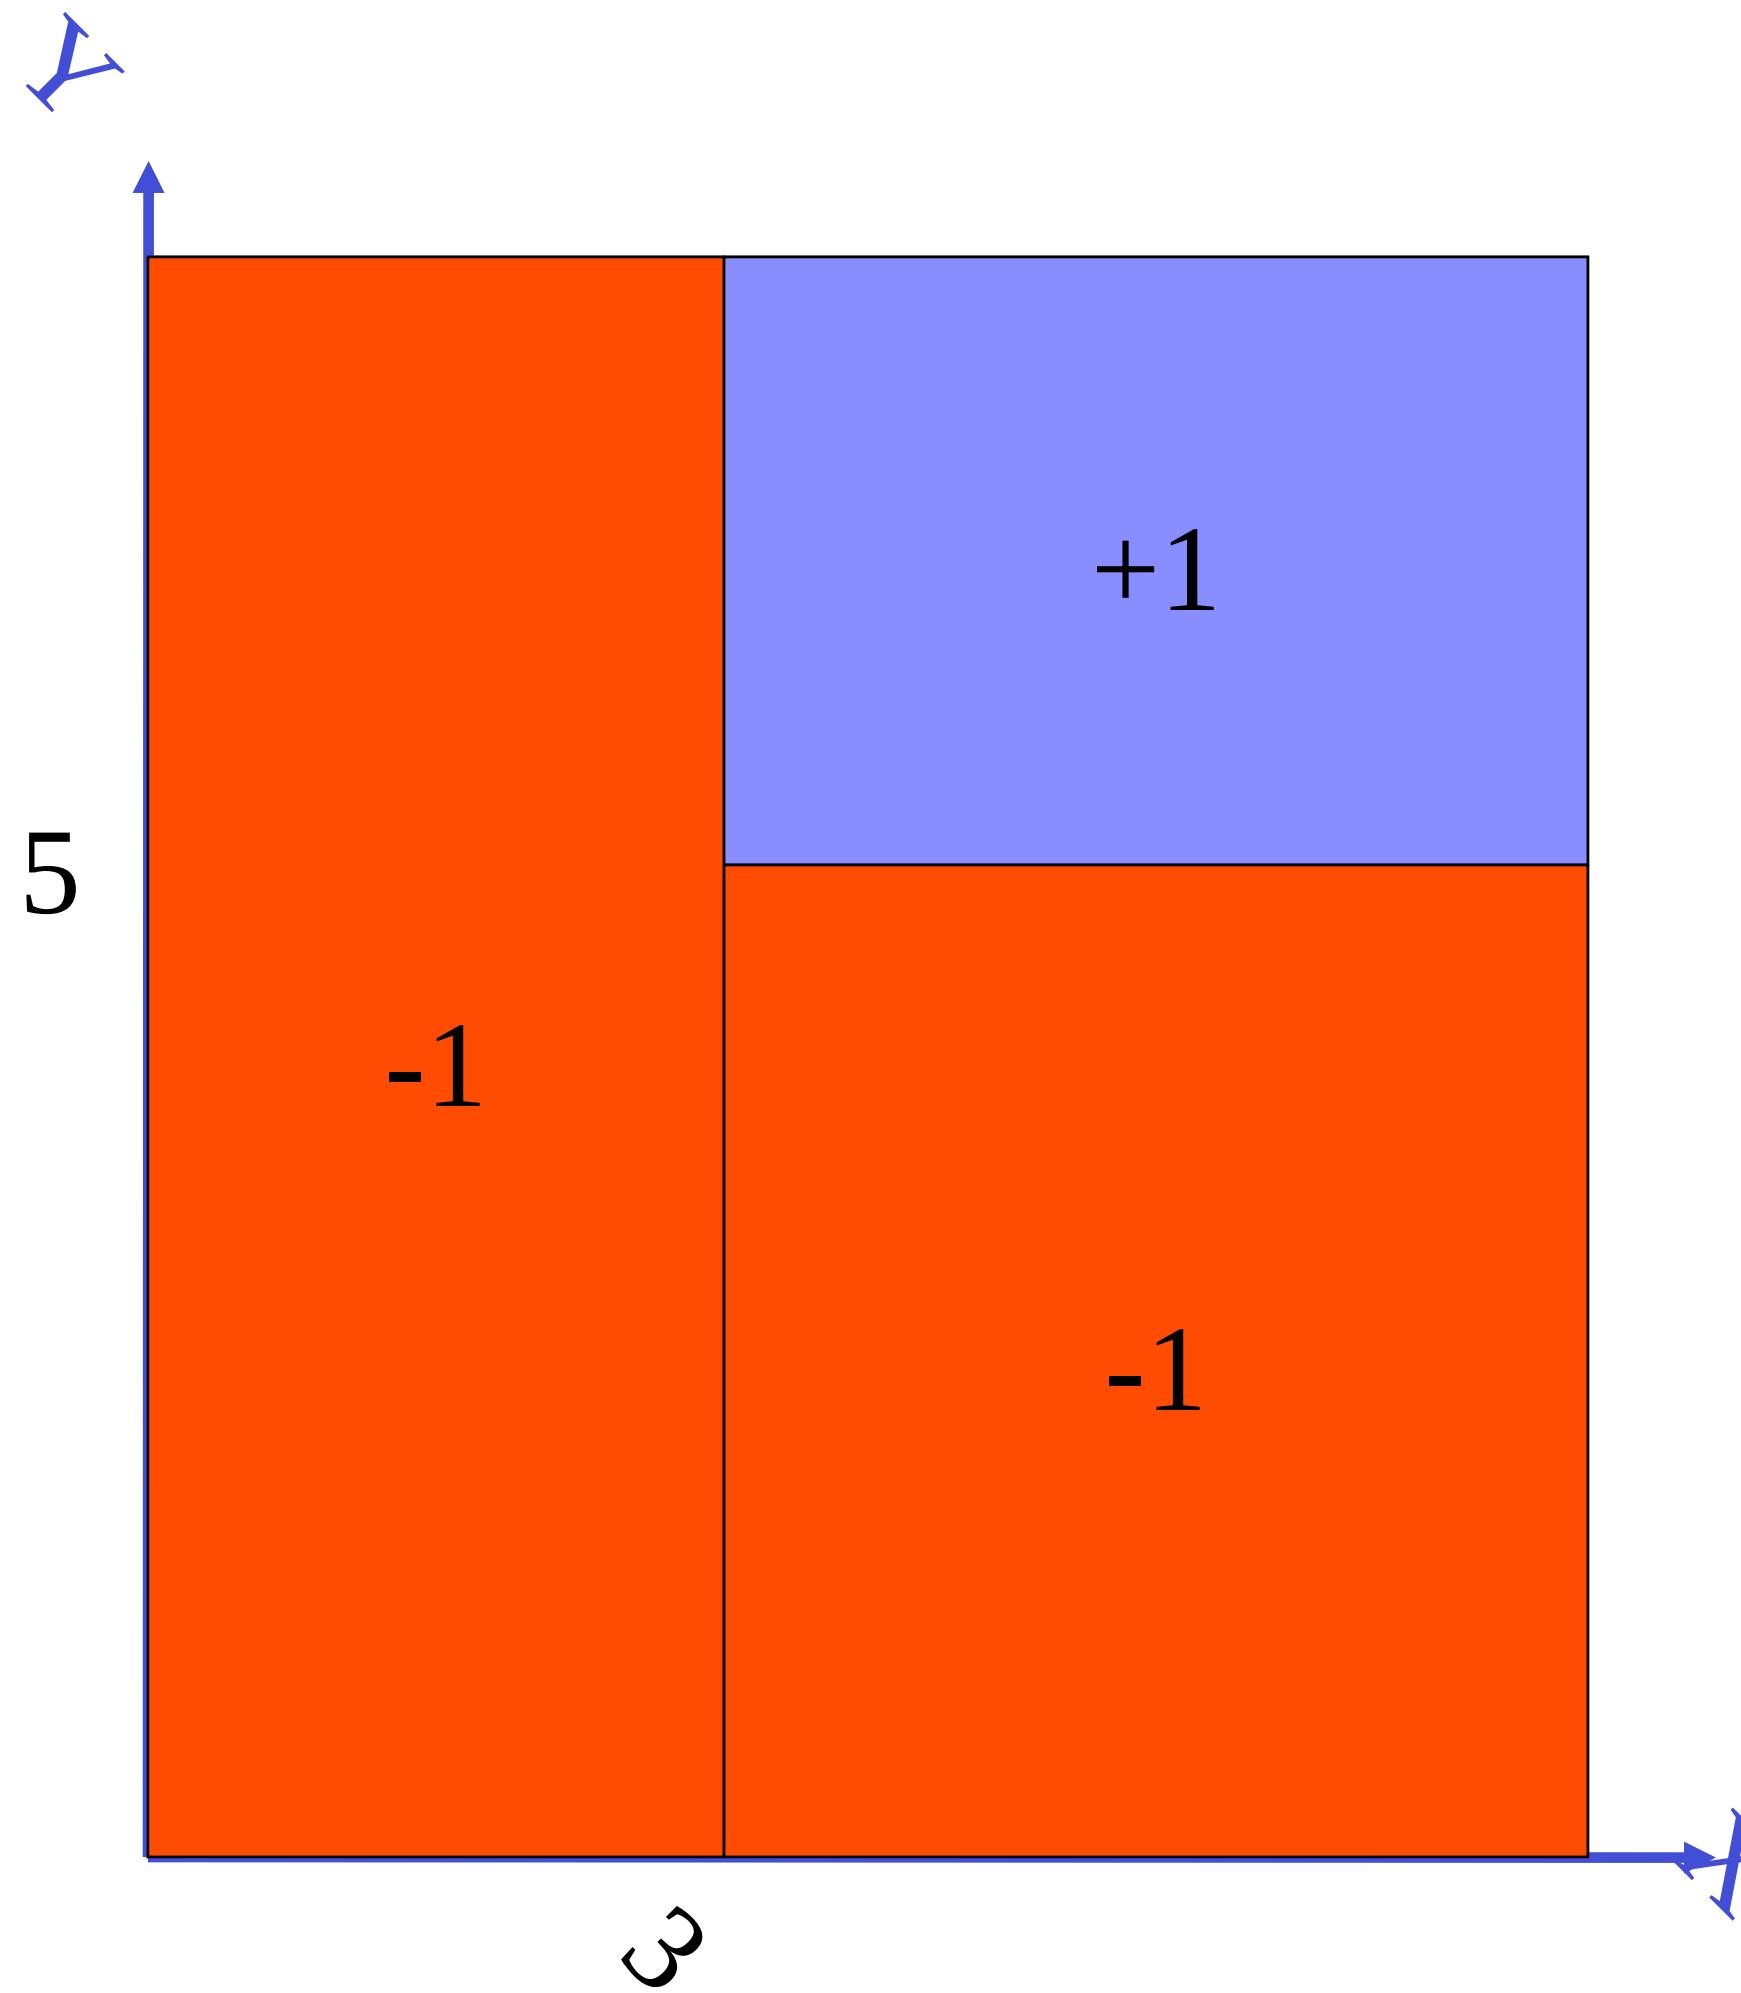
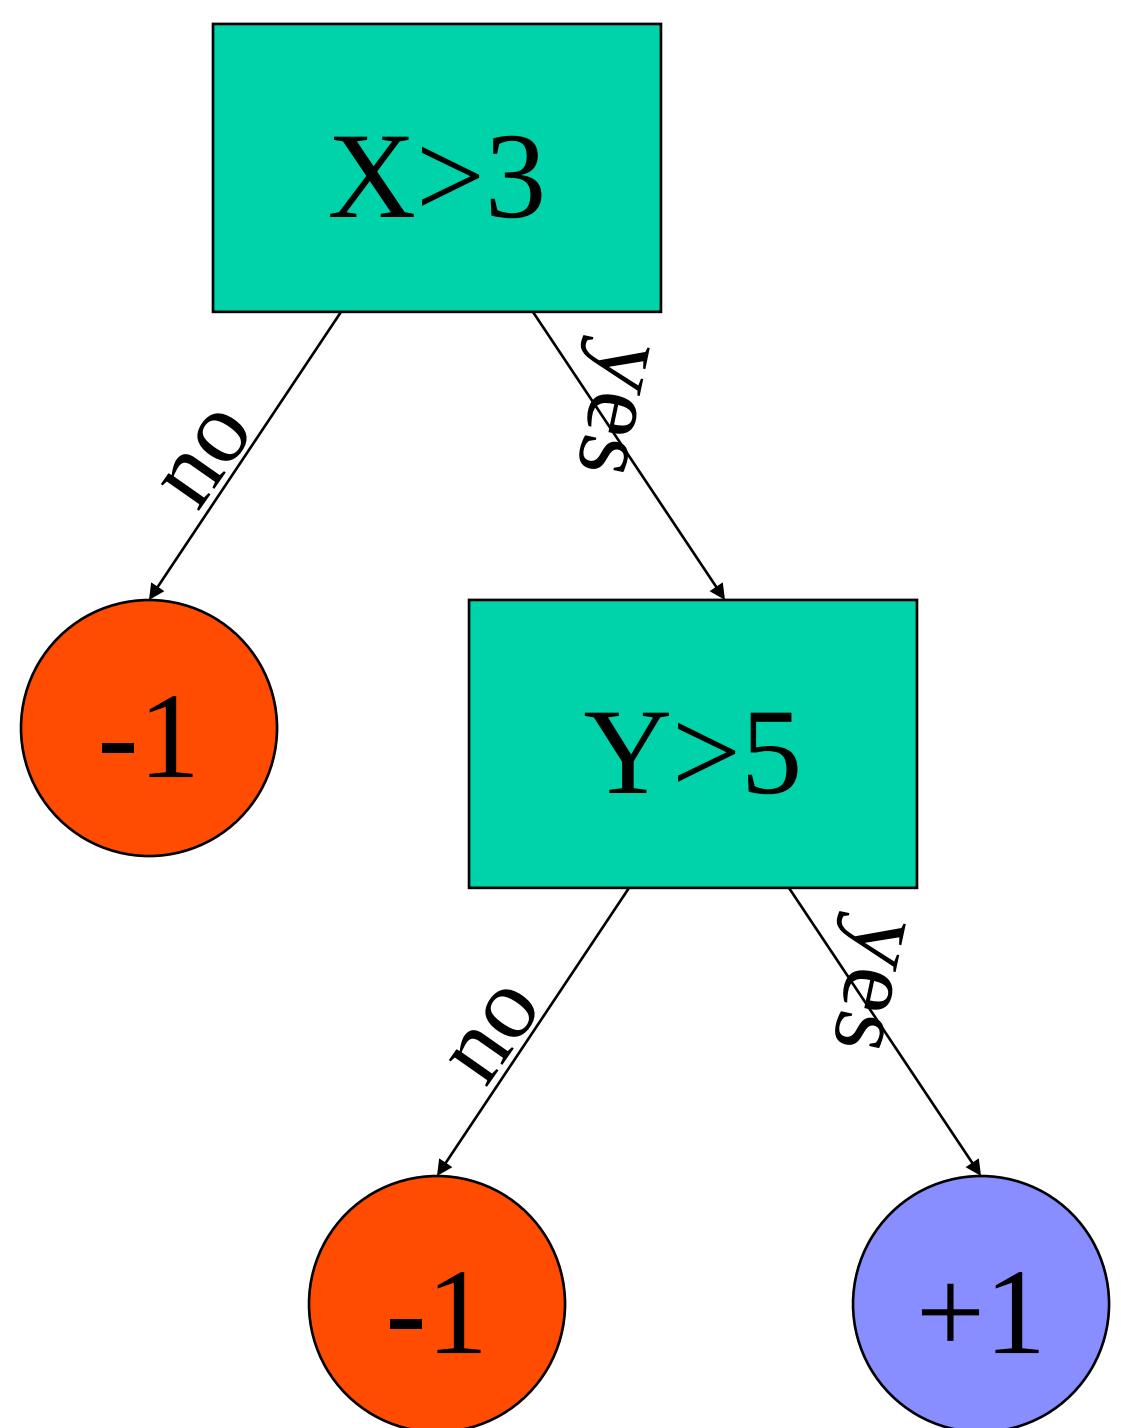
# Decision Trees / Continuous Features

Will bomb explode?

Mass	Temperature	explosion
1	100	-1
3.4	945	-1
10	32	-1
11.5	1202	+1



# Decision Trees



# Decision trees

- Popular because very flexible and easy to interpret.
- Learning a decision tree = finding a tree with small error on the training set.
  1. Start with the root node.
  2. At each step split one of the leaves
  3. Repeat until a termination criterion.

# Which node to split?

- We want the children to be more “pure” than the parent.
- Example:
  - Parent node is 50%+, 50%-.
  - Child nodes are (90%+,10%-),(10%+,90%-)
- How can we quantify improvement in purity?

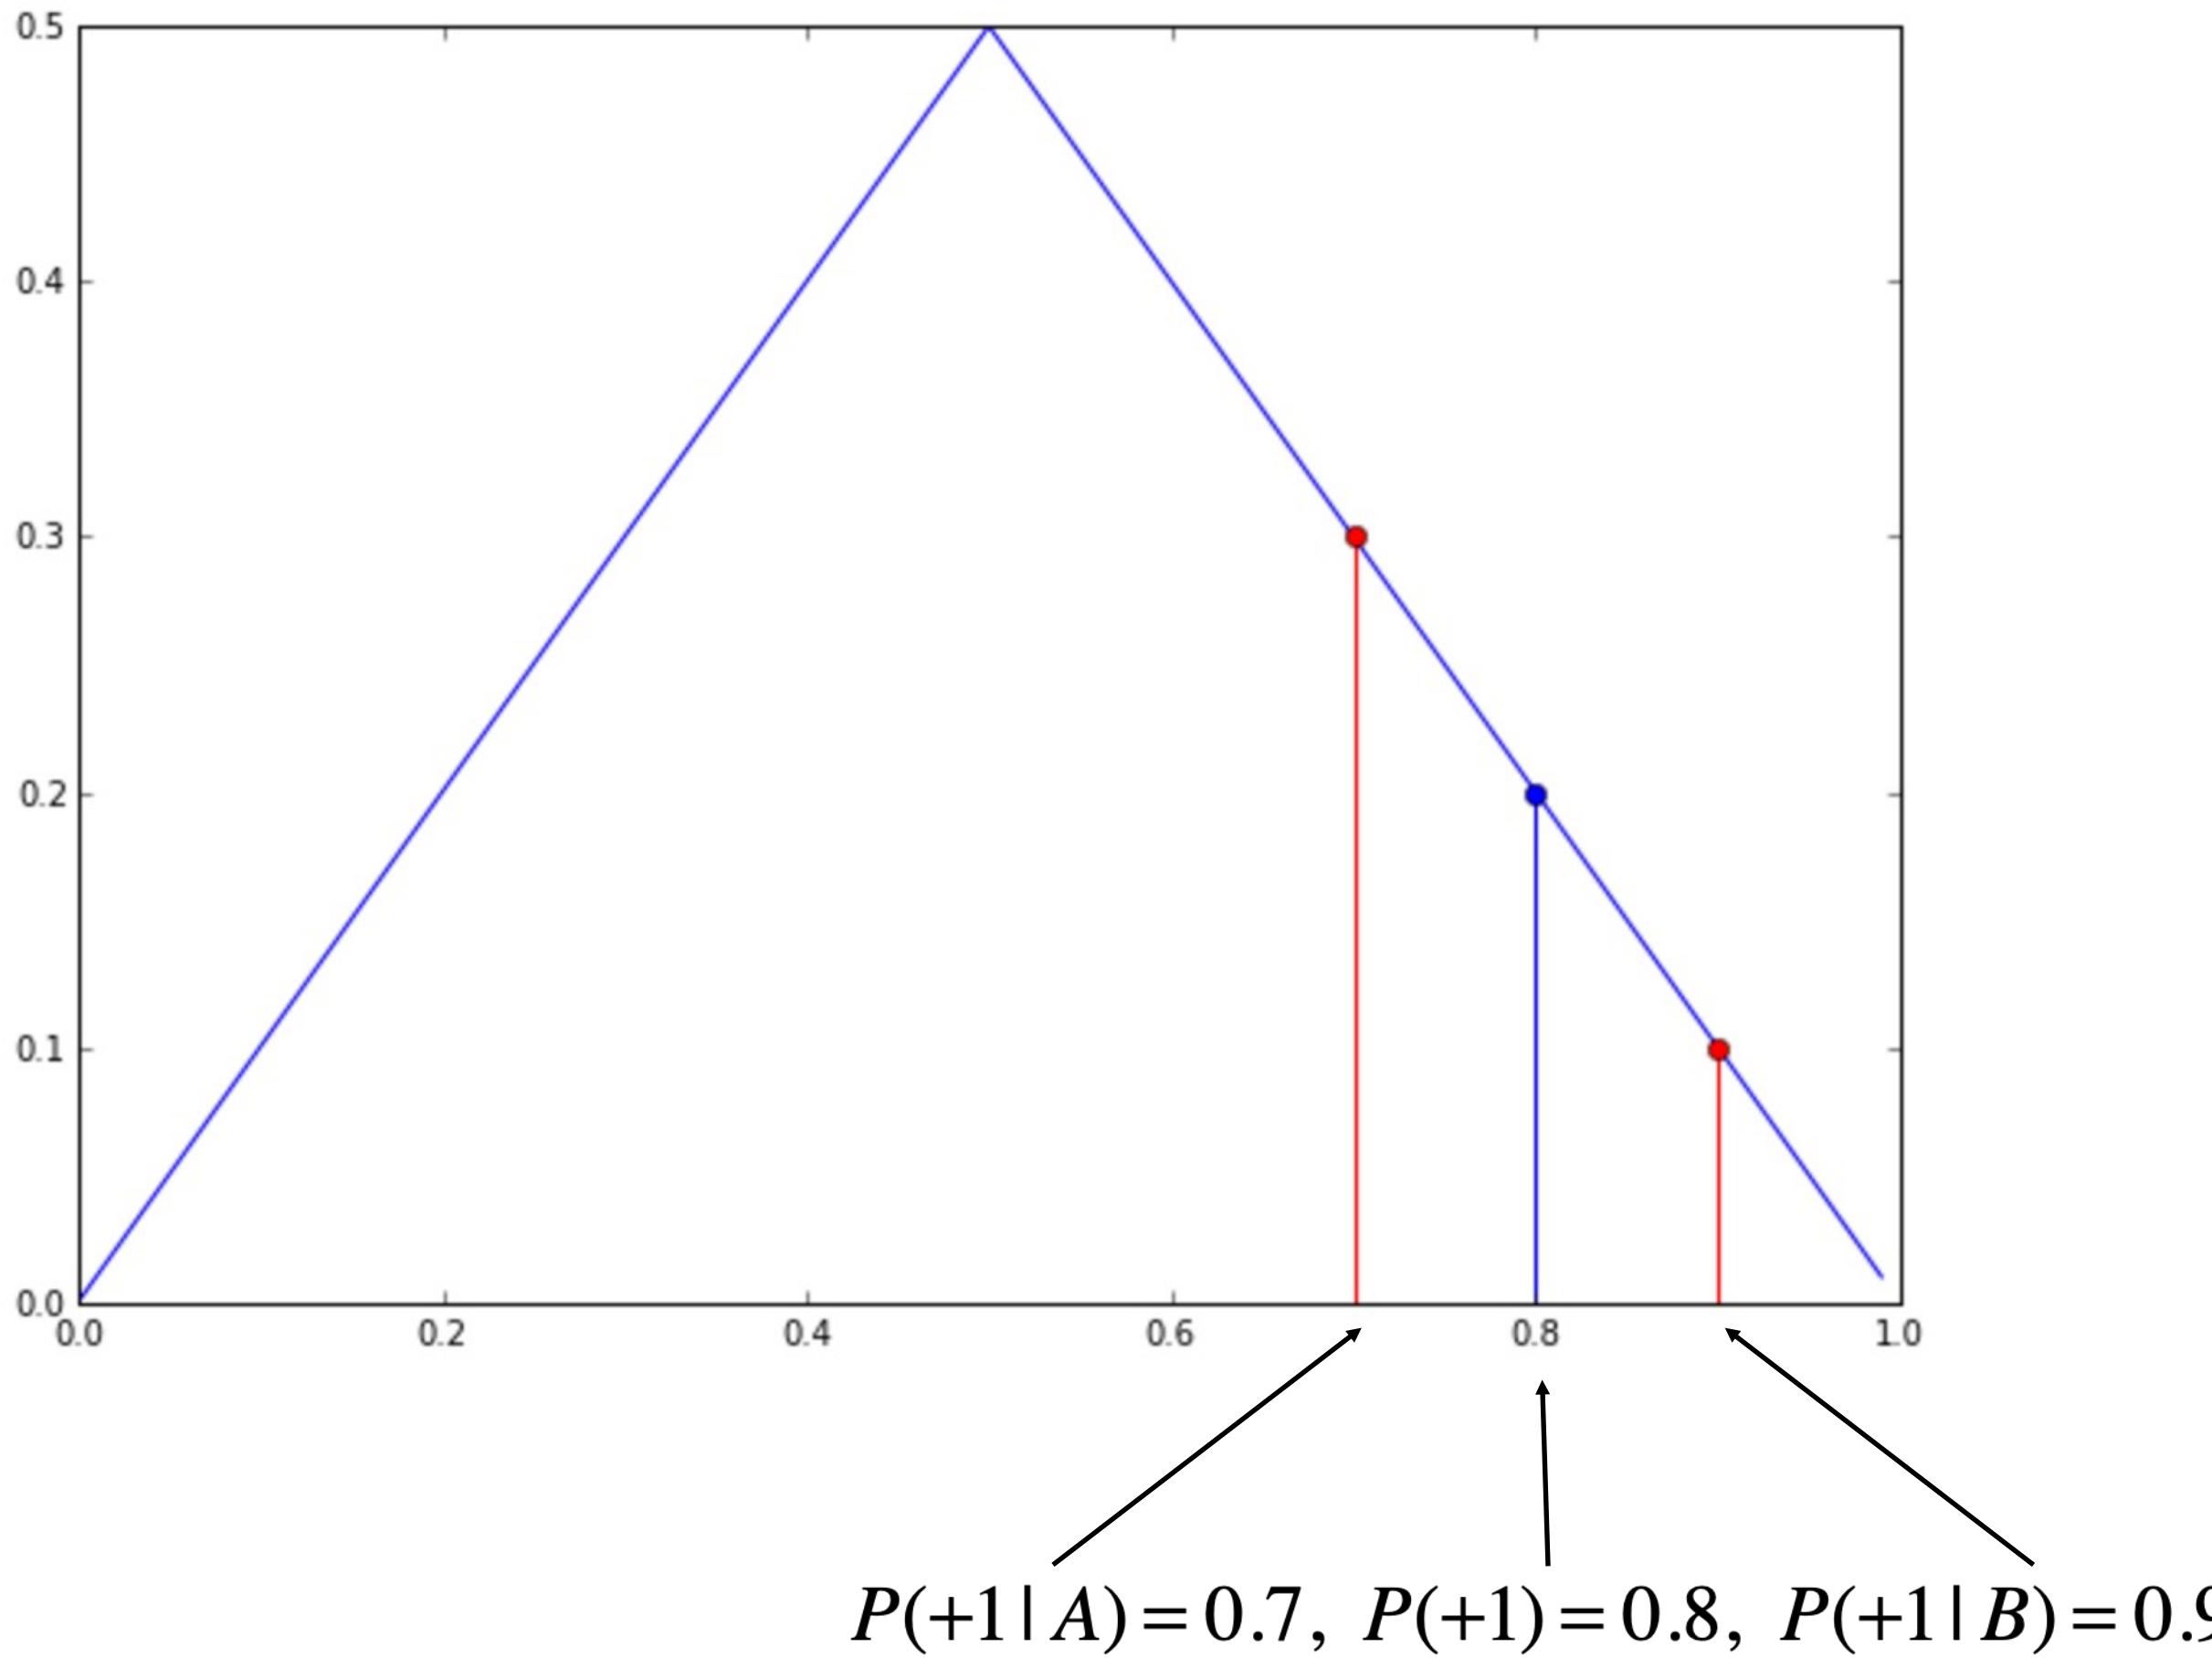
# Naive approach: minimize training error

- A good case
- Parent node is 55%+, 45%-.
  - Predict + always: Error = 45%
  - Child nodes are (90%+,10%),(10%+,90%)
    - Predict + on left, - on right: Error = 10%
- Clear Improvement

# Naive approach: minimize training error

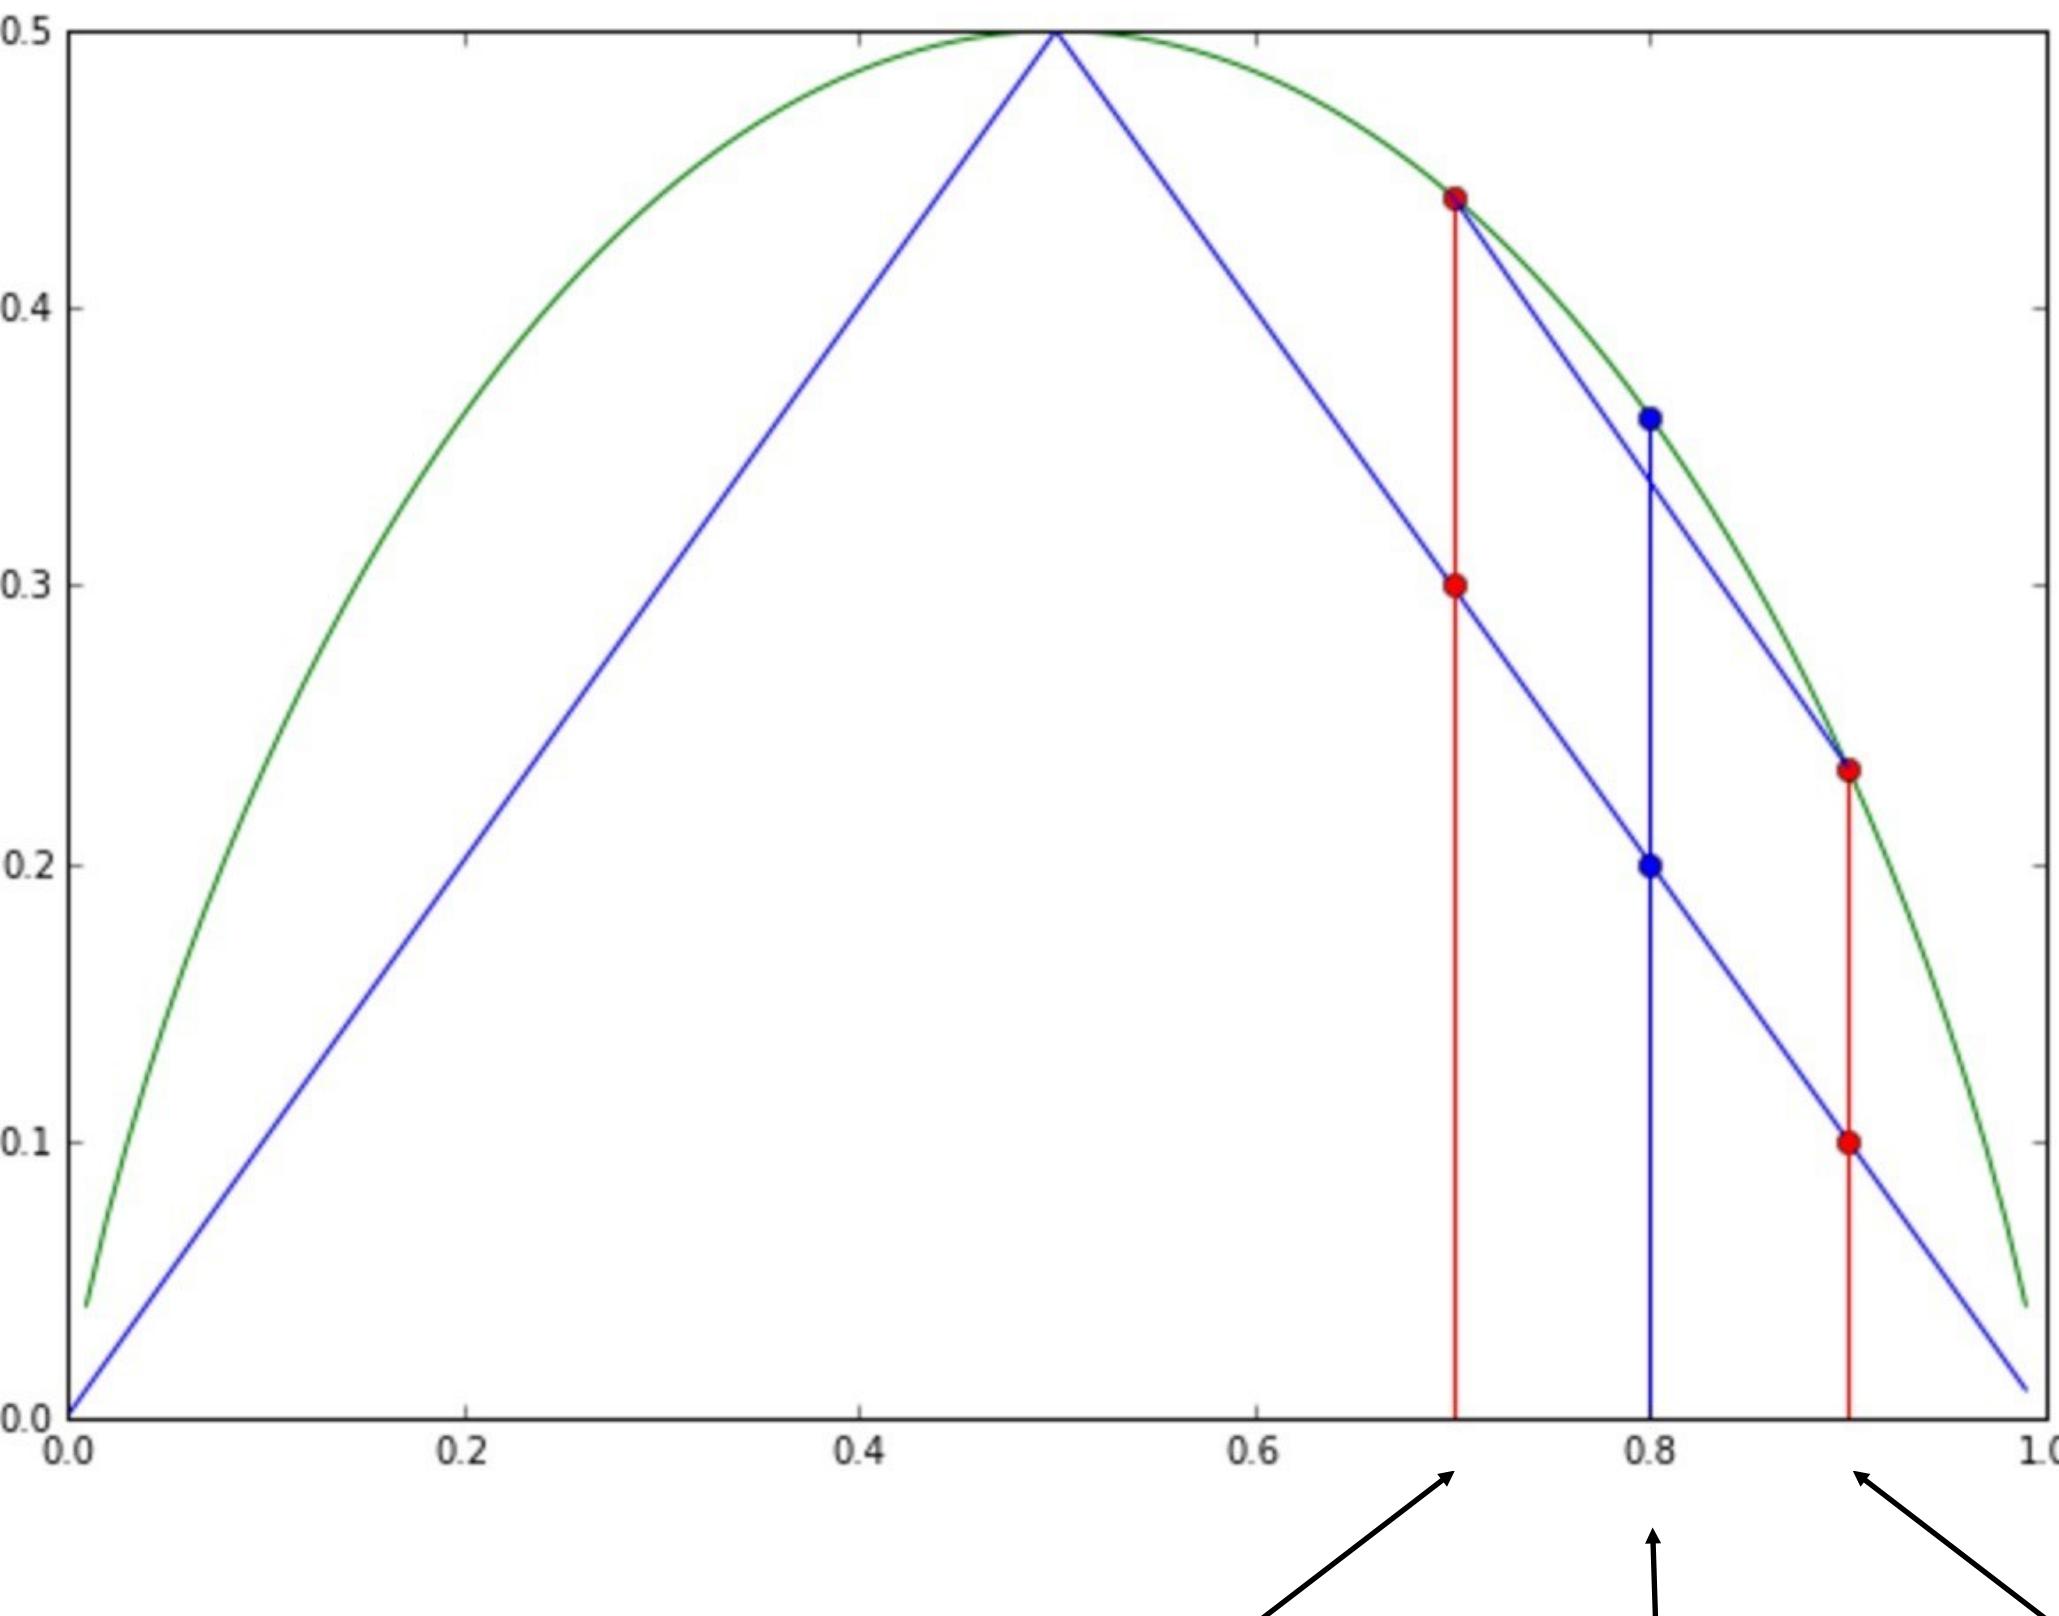
- **A bad case**
- Parent node is 80%+, 20%-.
  - Predict + always: Error = 20%
- Child nodes are (90%+,10%-),(70%+,30%-)
  - Predict + always: Error = 20%
- No improvement in training error!

# The problem with classification error (pictorially)



# Fixing the problem

instead of  $\text{err}(p) = \min(p, 1-p)$  use  $\frac{H(p)}{2} = -\frac{1}{2}(p \log_2 p + (1-p) \log_2(1-p))$



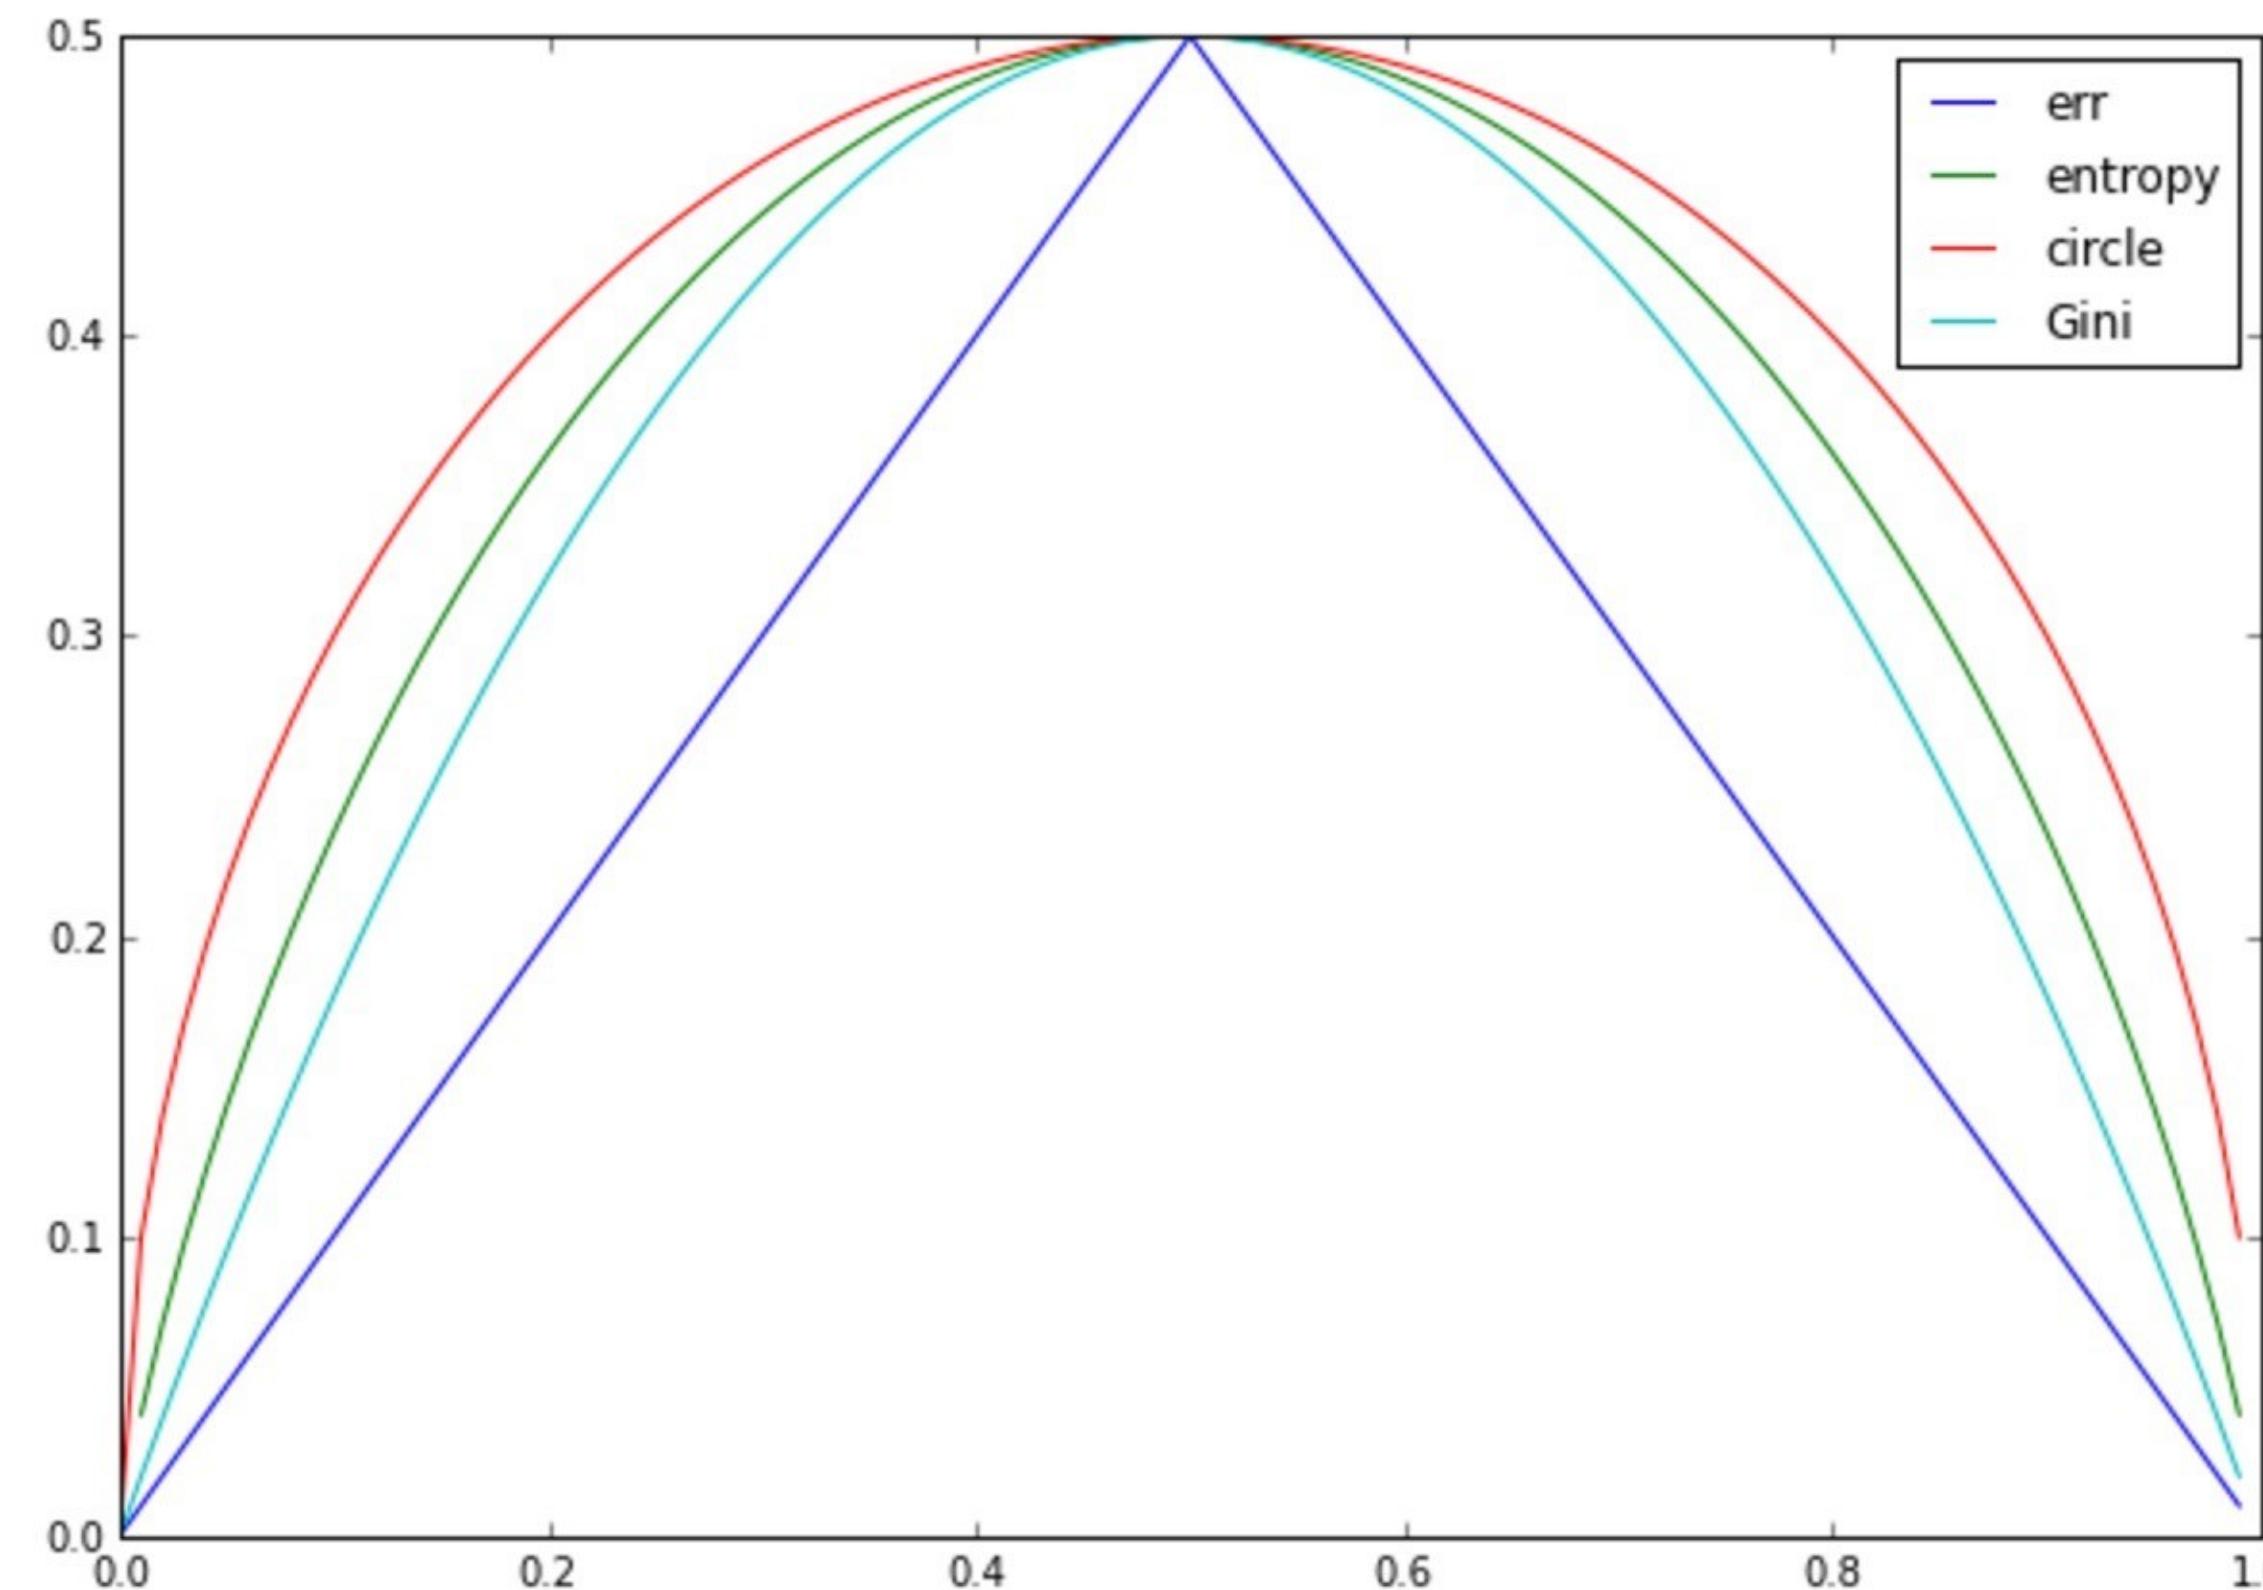
$$P(+1|A)=0.7, P(+1)=0.8, P(+1|B)=0.9$$

# Any strictly concave function can be used

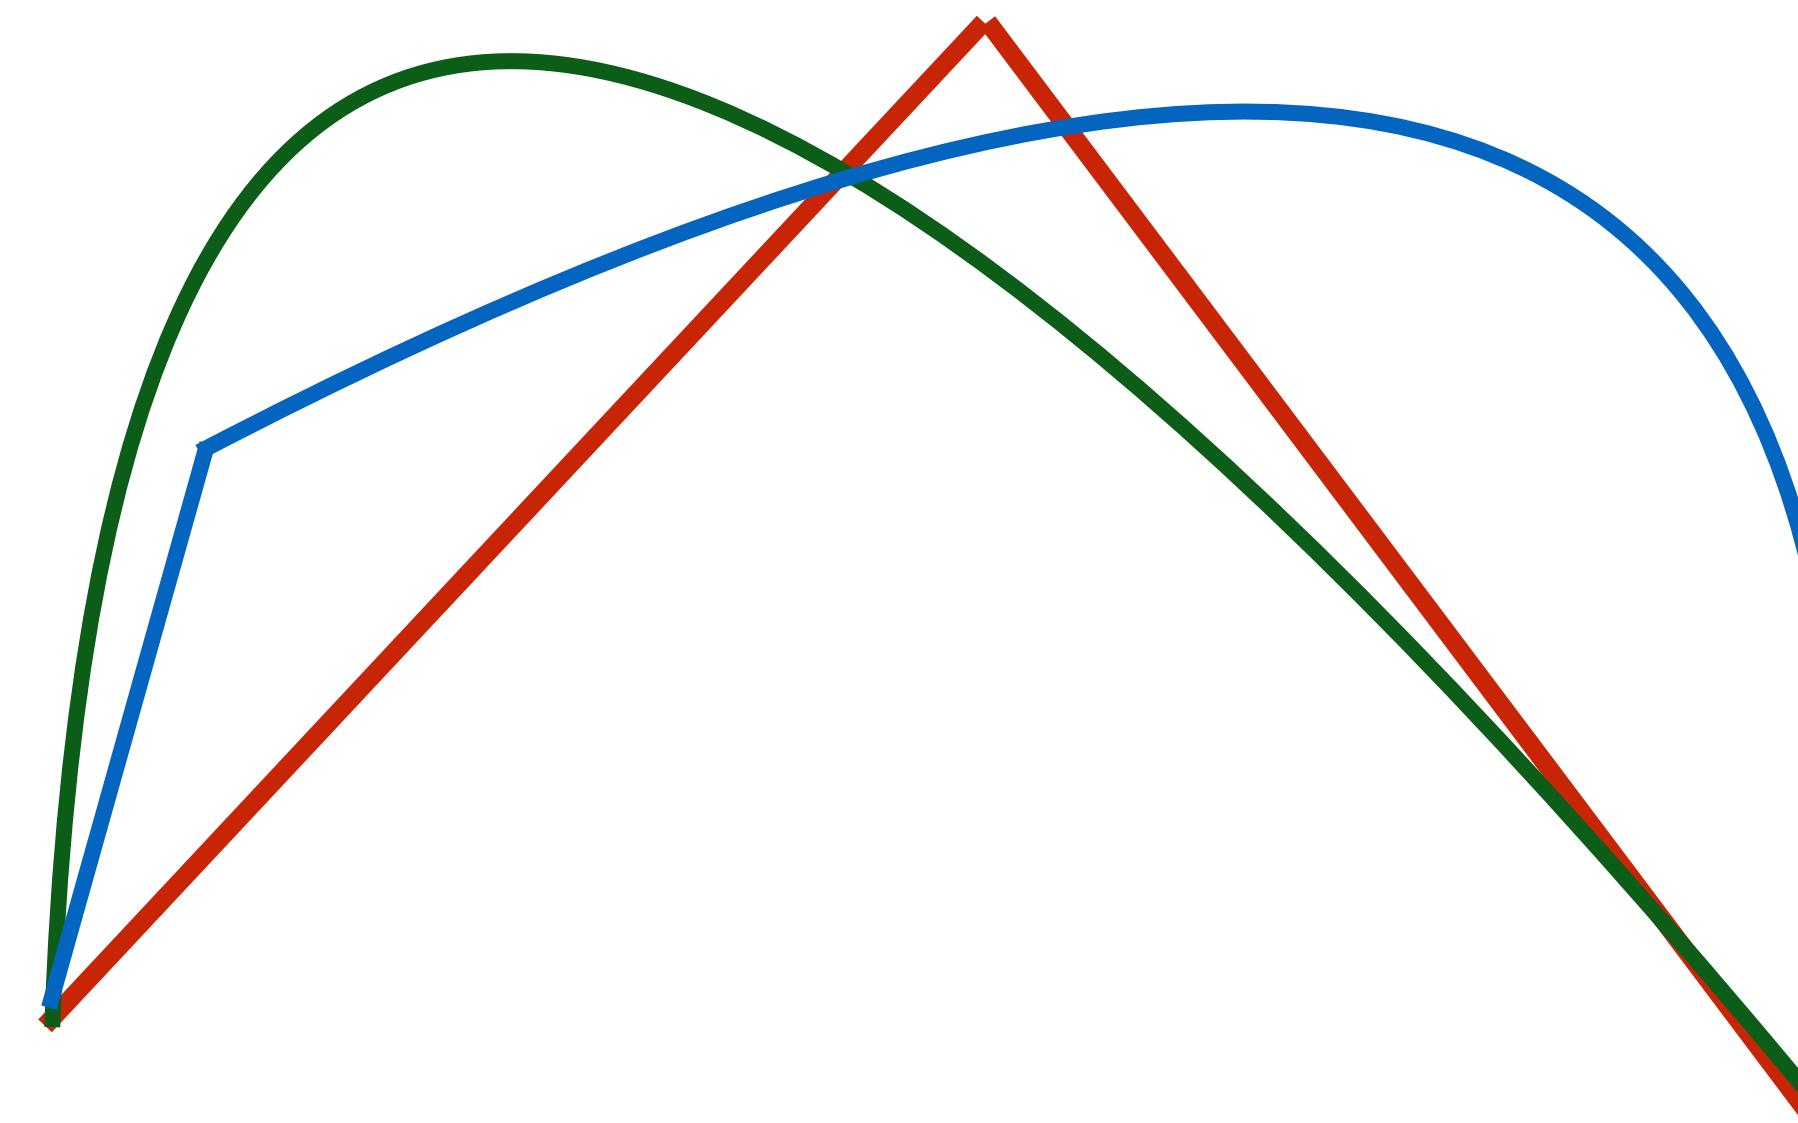
$$H(P) = p \log p + (1-p) \log(1-p)$$

$$\text{Circle}(p) = \sqrt{1/4 - (p - 1/2)^2}$$

$$\text{Gini}(p) = p(1-p)$$



Which of the following is  
strictly concave?



# Decision tree learning algorithm

- Learning a decision tree = finding a tree with small error on the training set.
  1. Start with the root node.
  2. At each step split one of the leaves
  3. Repeat until a termination criterion.
  4. Next, How do we search for a splitting rule.

# The splitting step

- Given: current tree.
- For each leaf and each feature,
  - find all possible splitting rules (finite because data is finite).
  - compute reduction in entropy
- find leaf X feature X split rule the minimizes entropy.
- Add selected rule to split selected leaf.

# Enumerating splitting rules

- If feature has a fixed, small, number of values.  
then either:
  - Split on all values (Location is beach/prison/ski-slope)
  - or Split on equality to one value (location = beach)
- If feature is continuous (temperature) then either:
  - Sort records by feature value and search for best split.
  - or split on percentiles: 1%, 2%, ..., 99%

# Splitting on percentiles

- Suppose data is in an RDD with 100 million examples.
- sorting according to each feature value is very expensive.
- Instead: use `Sample(false,0.00001).collect()` to get a sample of about 10,000 examples.
- sort the sample (small, sort in head node).
- pick examples at location 100,200,... as boundaries. Call those feature values  $T_1, T_2, T_3, \dots, T_{99}$
- Broadcast boundaries to all partitions.
- Each partition computes it's contribution to  $P(+1 | T_i \leq f \leq T_{i+1})$

# Summary

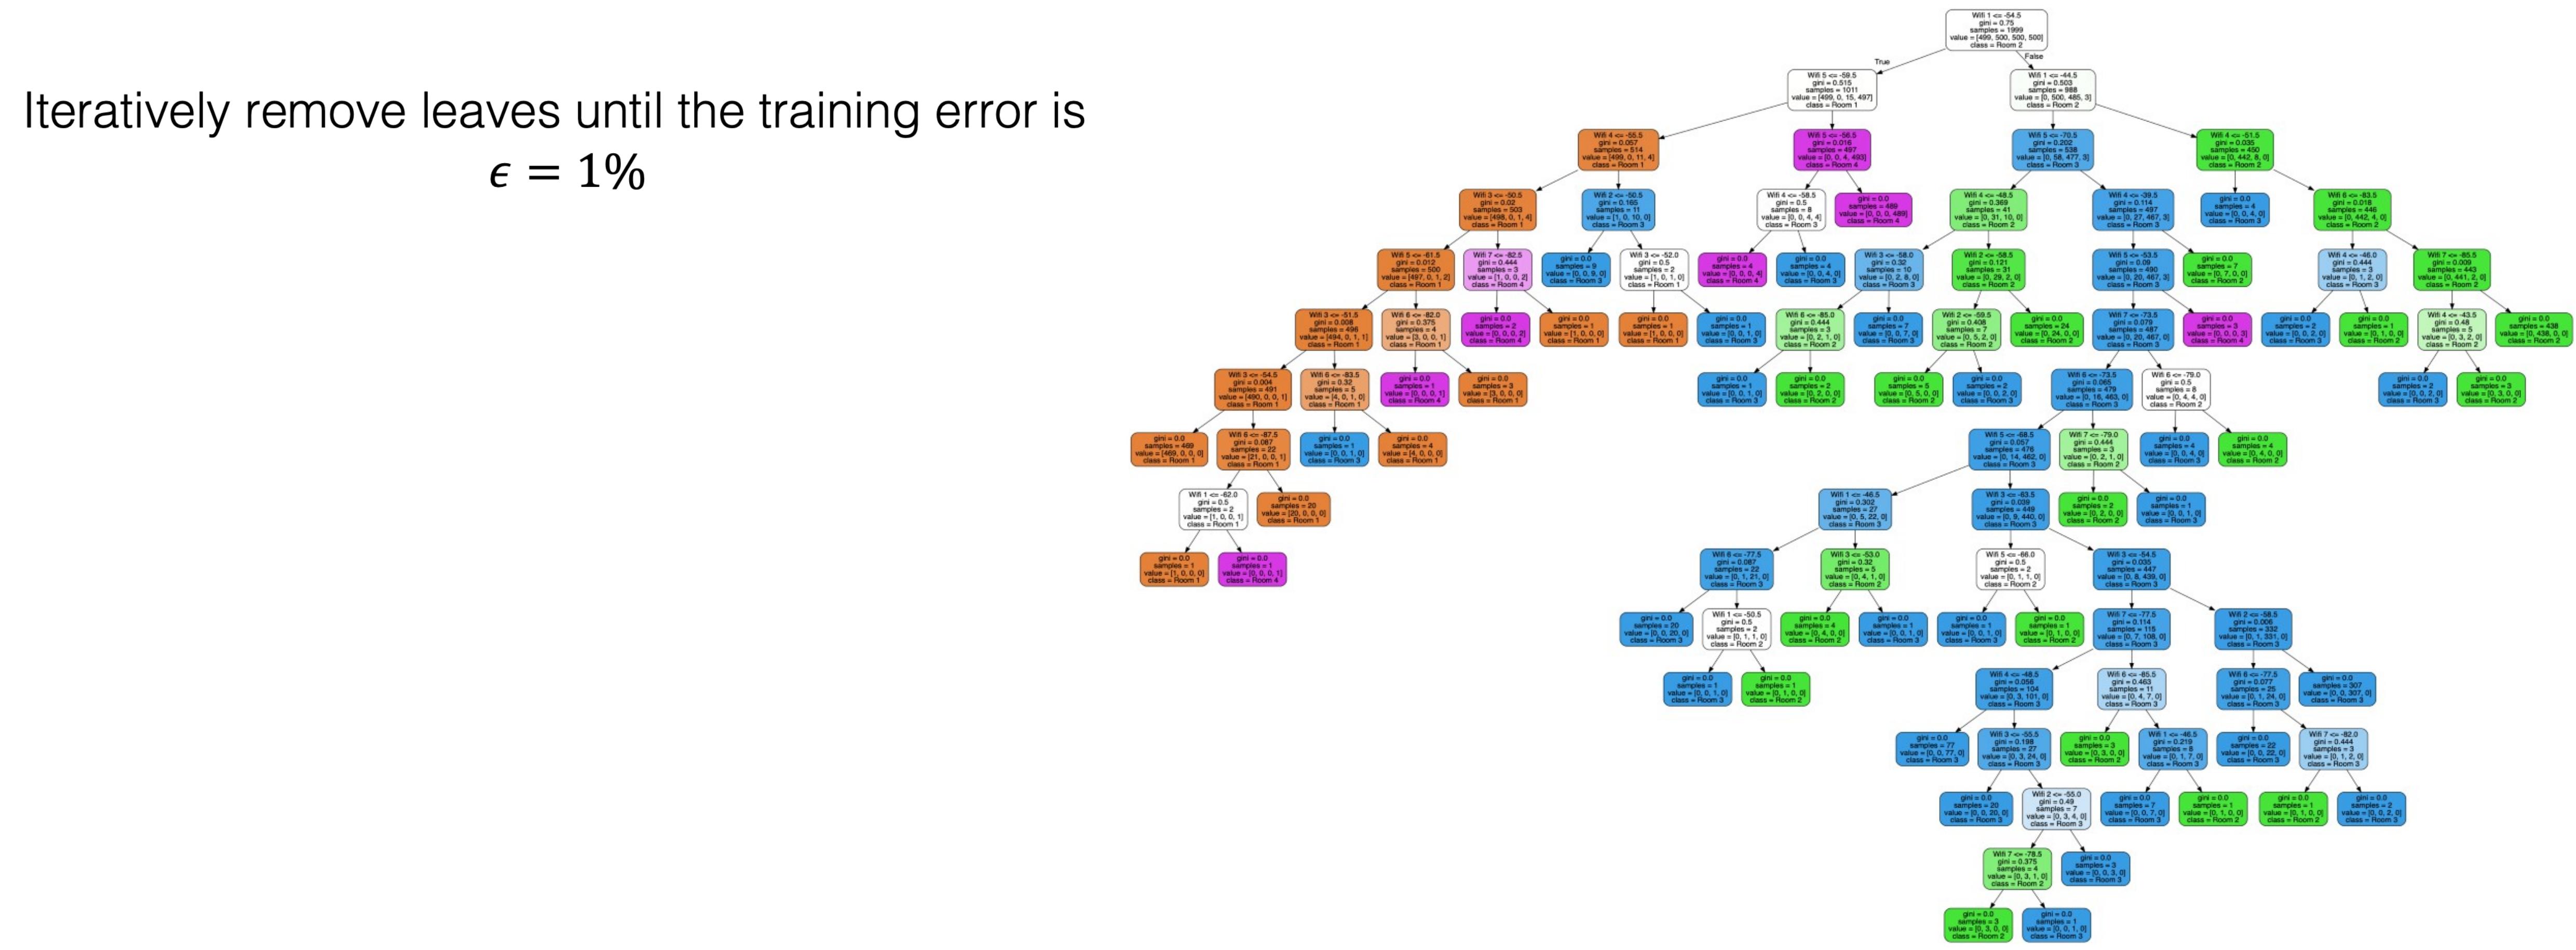
- The splitting criteria is a strictly concave function that bounds the training error.
- The search for a splitting rule tests all possible splits of current tree.
- When the feature is continuous and the data large - split on percentiles instead of splitting on each example.
- Next, Performance on the test set and reducing over-fitting .

# Trees are unstable

- Trees are very flexible.
- A “fully grown” tree is one where all leaves are “pure”, i.e. each leaf contains all + labeled examples or all - labeled examples.
- A fully grown tree has training error zero.
- If the tree is large and the data is limited, the test error of the tree is likely to be high = the tree overfits the data.
- Statisticians say that trees are “high variance” or “unstable”

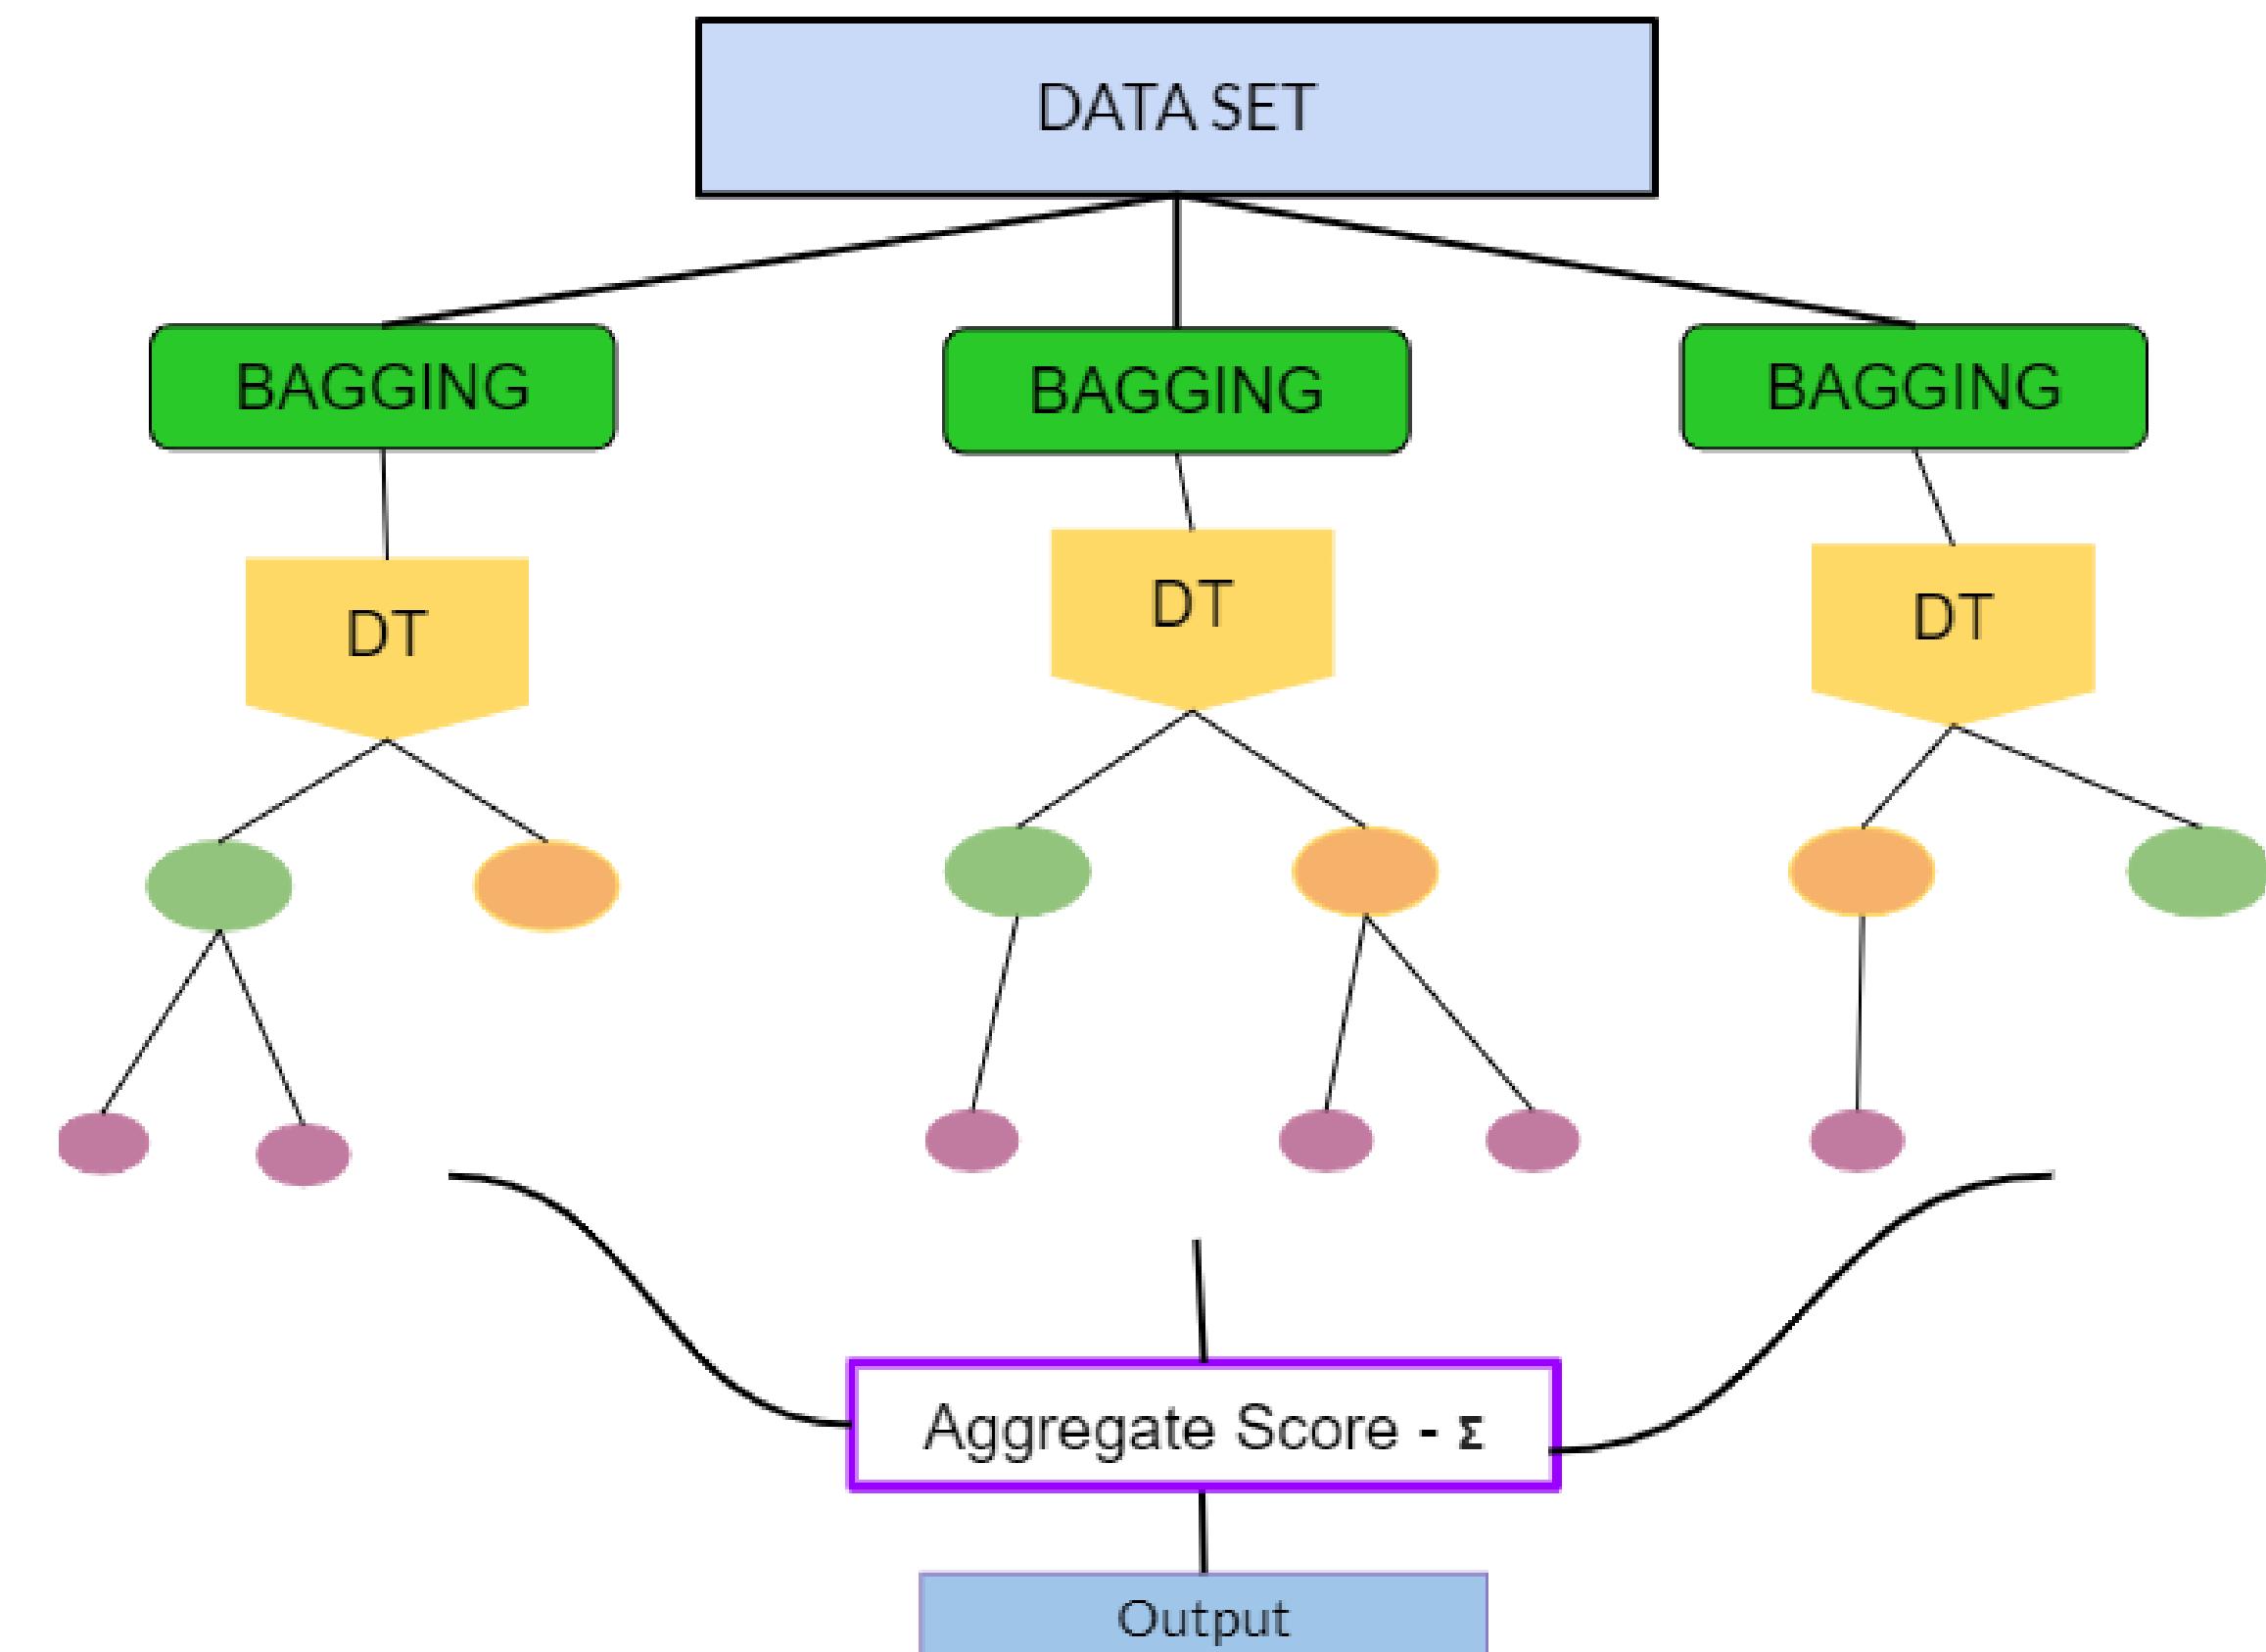
# Pruning Trees

- One way to make trees more stable is to prune.
- Start by building a tree down to training error of zero.
- Iteratively remove leaves until the training error is  $\epsilon = 1\%$



# Bagging Trees

- Bagging, invented by Leo Breiman in the 90s, is a different way to reduce the variance of trees.
- Instead of pruning the tree, we generate many trees, using randomly selected subsets of the training data.
- We predict using the majority vote over the trees.
- Related to RANDOM FORESTS and to Boosted Trees/XGBoost. Will describe next.



# Summary

- Decision trees are a simple and intuitive representation.
- Fully grown trees over-fit the data.
- Two ways to reduce over-fitting:
  - Pruning: remove leaves to make the tree smaller.
  - Bagging, Boosting: take the majority vote over many trees.

# Are methods based on decision tree still relevant at the age of Deep Neural Networks?

## Kaggle Survey of most popular learning algorithms

(2020)

